

Low-level Vision Processing Algorithms

Speaker: Ito, Dang

Supporter: Ishii, Toyama and Y. Murakami

Adaptive Systems Lab
The University of Aizu

- Introduction
 - What is Vision Processing?
- Basic Knowledge
 - Image formation
 - Transformation
- Low level Algorithm
 - Filtering
 - Edge/Border Detection
 - Conner Detection
- Conclusion

- Introduction
 - What is Vision Processing?
- Basic Knowledge
 - Image formation
 - Transformation
- Low level Algorithm
 - Filtering
 - Edge/Border Detection
 - Conner Detection
- Conclusion

- Make computers understand images and video:
 - Images to Models



What kind of scene?

Where are the cars?

How far is the building?

...

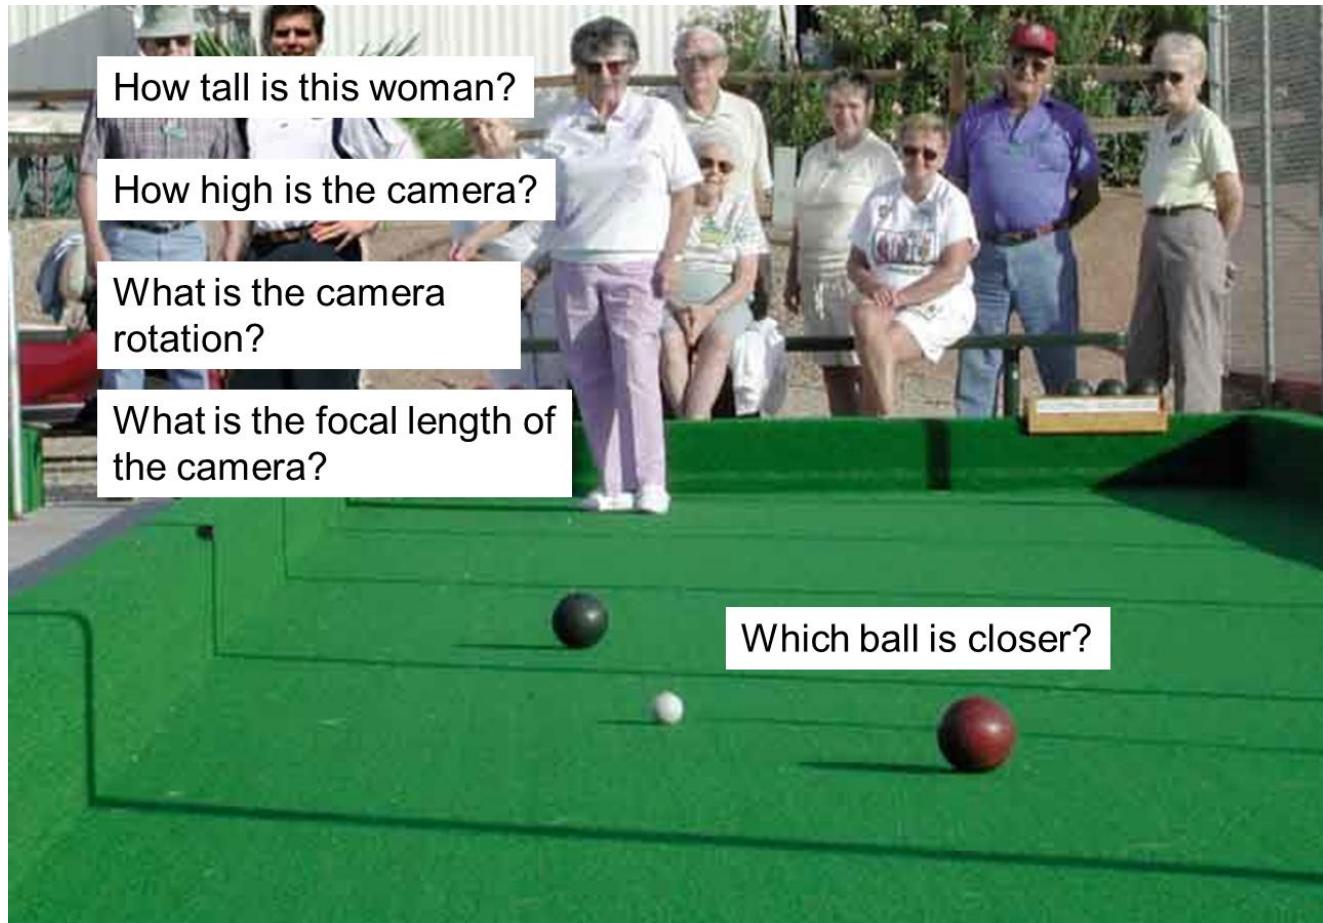
Slide credit James Hays



Application of Vision Processing

- Optical Character Recognition
- Face detection
- Object recognition
- Shape/motion capture
- Medical Imaging

- Human limitations
 - measurement accuracy, darkness, etc...

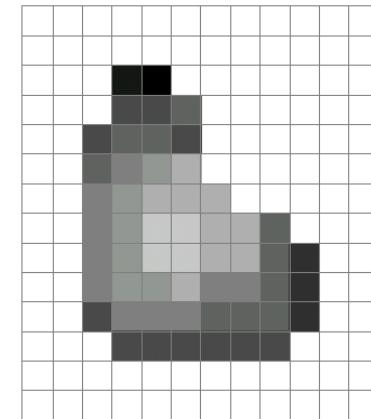
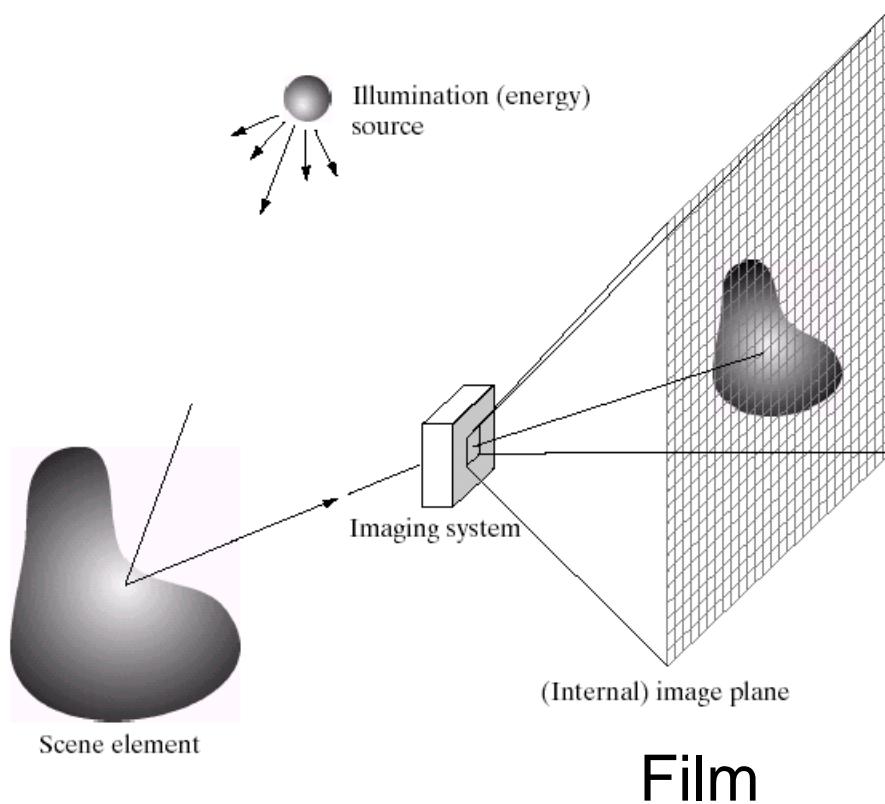


- Introduction
 - What is Vision Processing?
- Basic Knowledge
 - Image formation
 - Transformation
- Low level Algorithms
 - Filtering
 - Border Detection
 - Edge Detection
- Conclusion

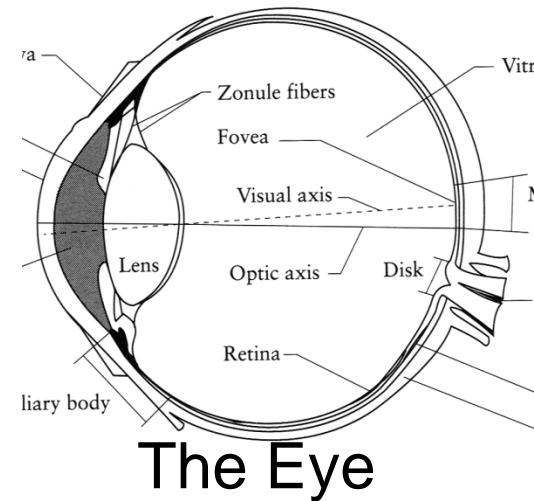


Basic Knowledge

- Image and Color.
- Camera, Lens.
- Color representation.
- Fourier Transform.



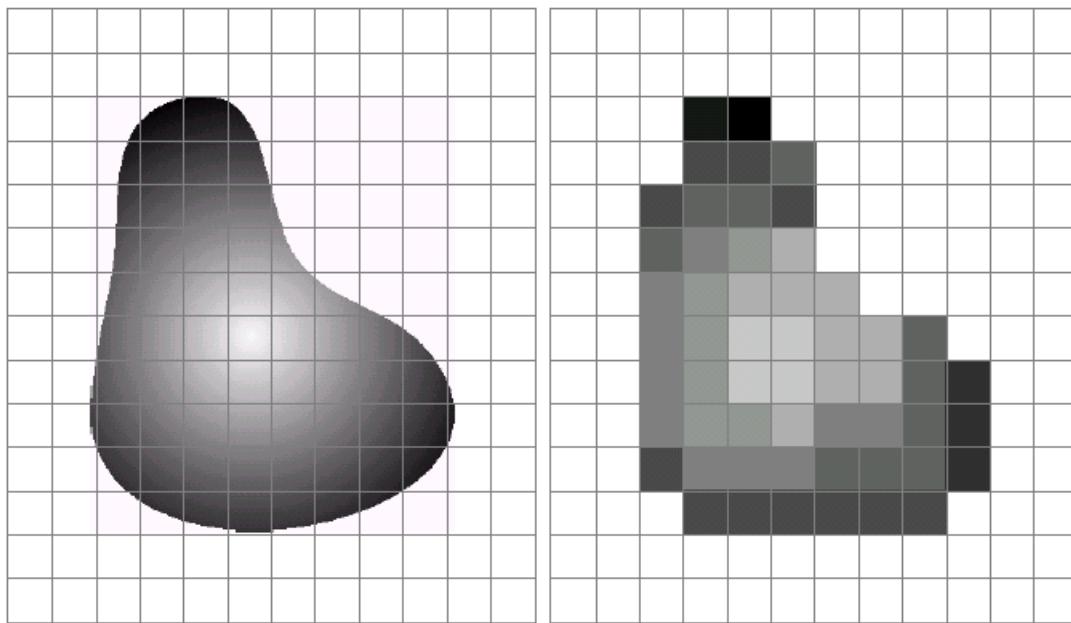
Digital Camera



The Eye

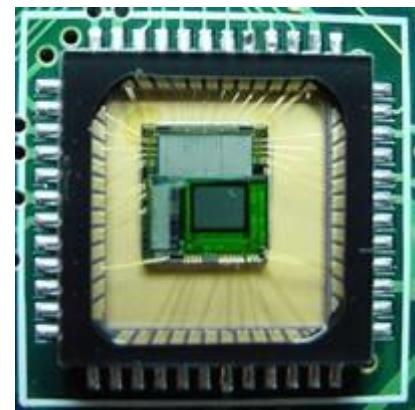


- A digital camera replaces film with a sensor array
 - Each cell in the array is light-sensitive diode that converts photons to electrons
 - Two common types
 - Charge Coupled Device (CCD)
 - CMOS
 - <http://electronics.howstuffworks.com/digital-camera.htm>

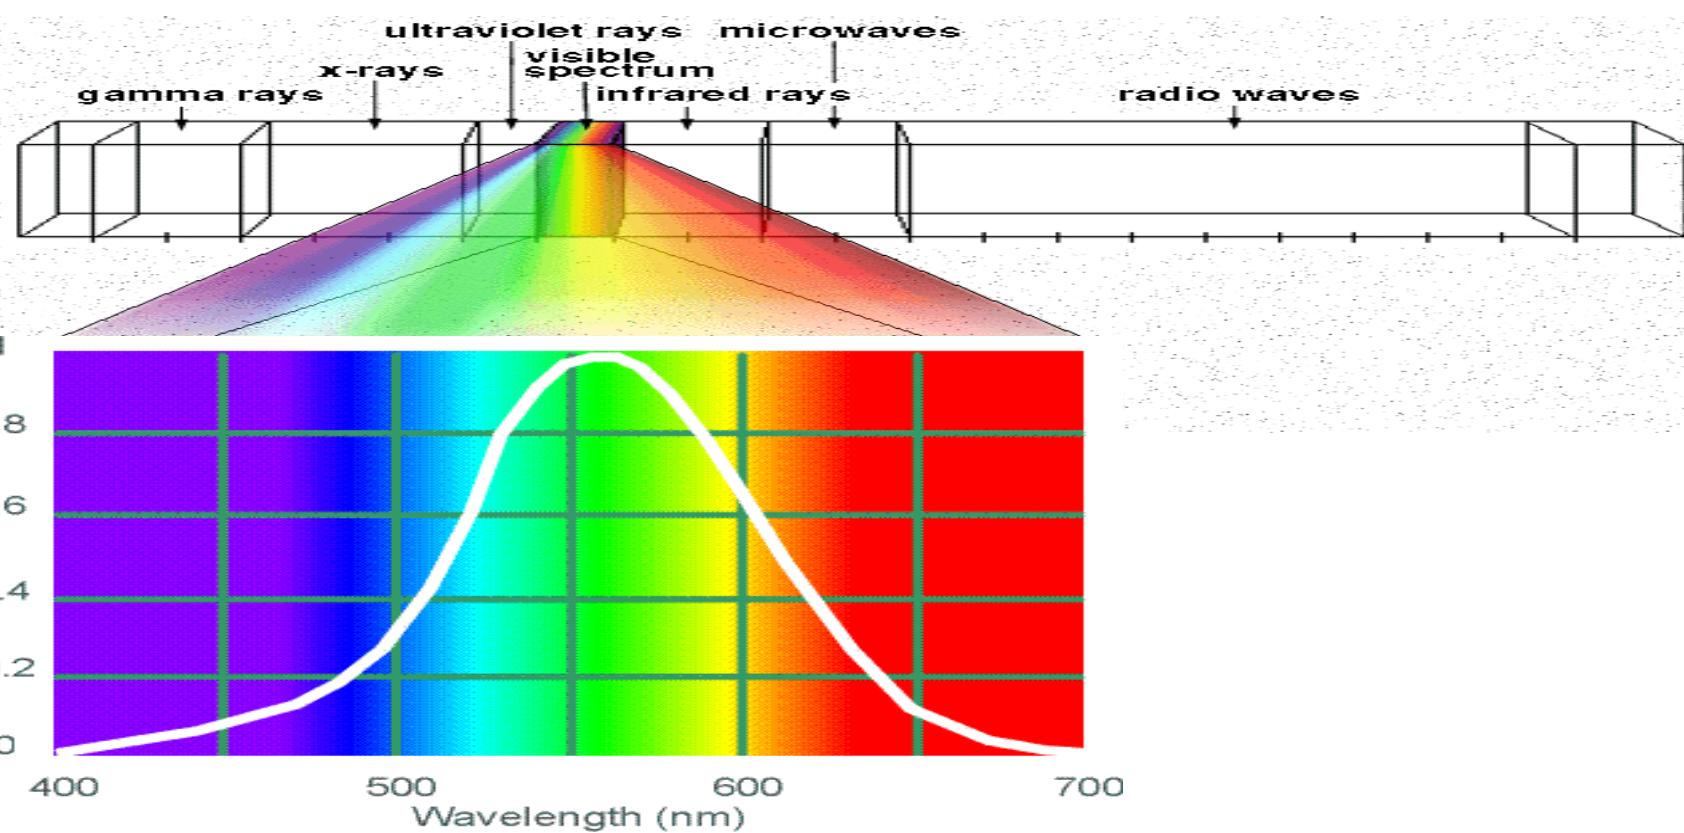


a b

FIGURE 2.17 (a) Continuos image projected onto a sensor array. (b) Result of image sampling and quantization.

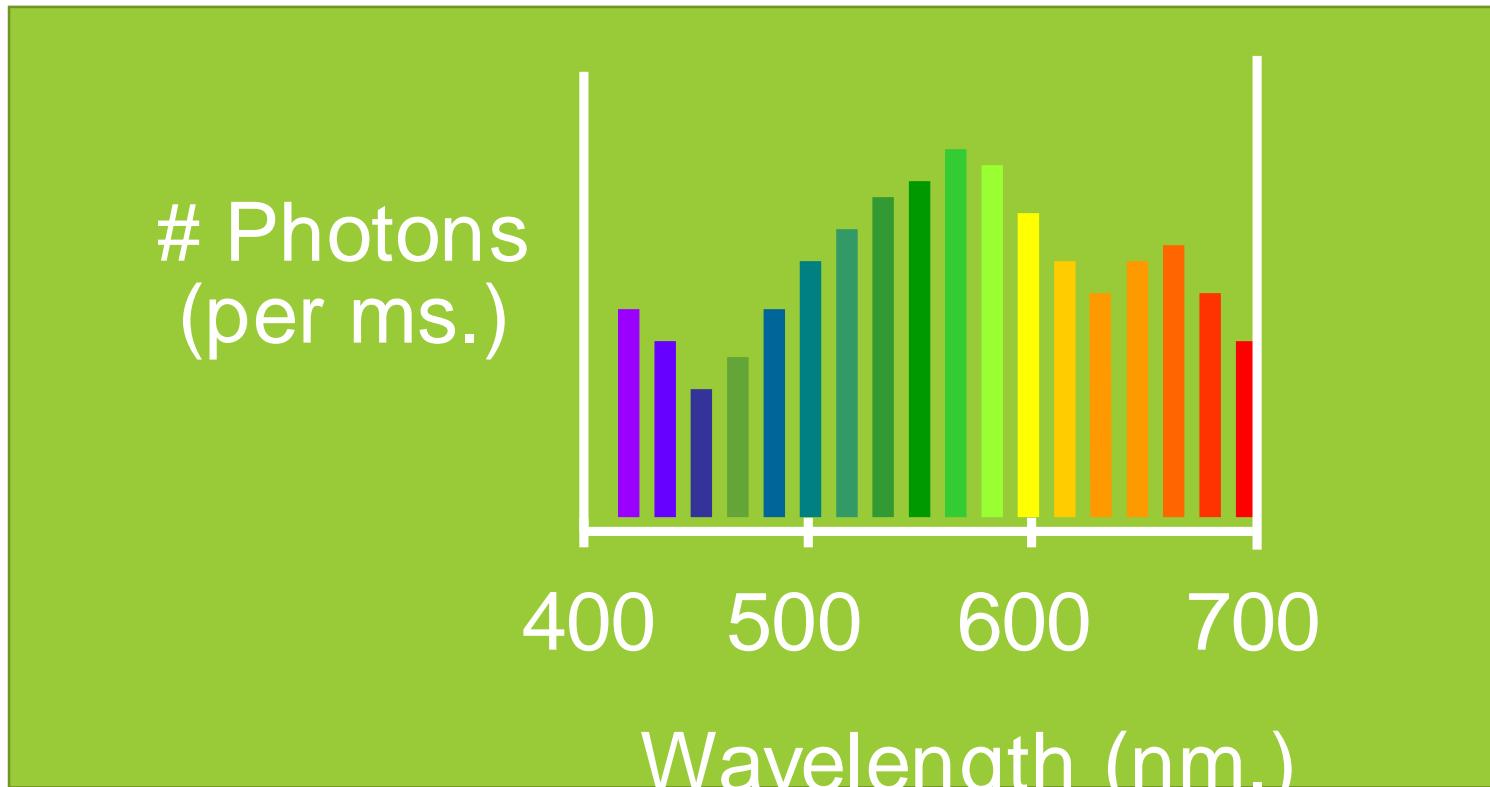


CMOS sensor



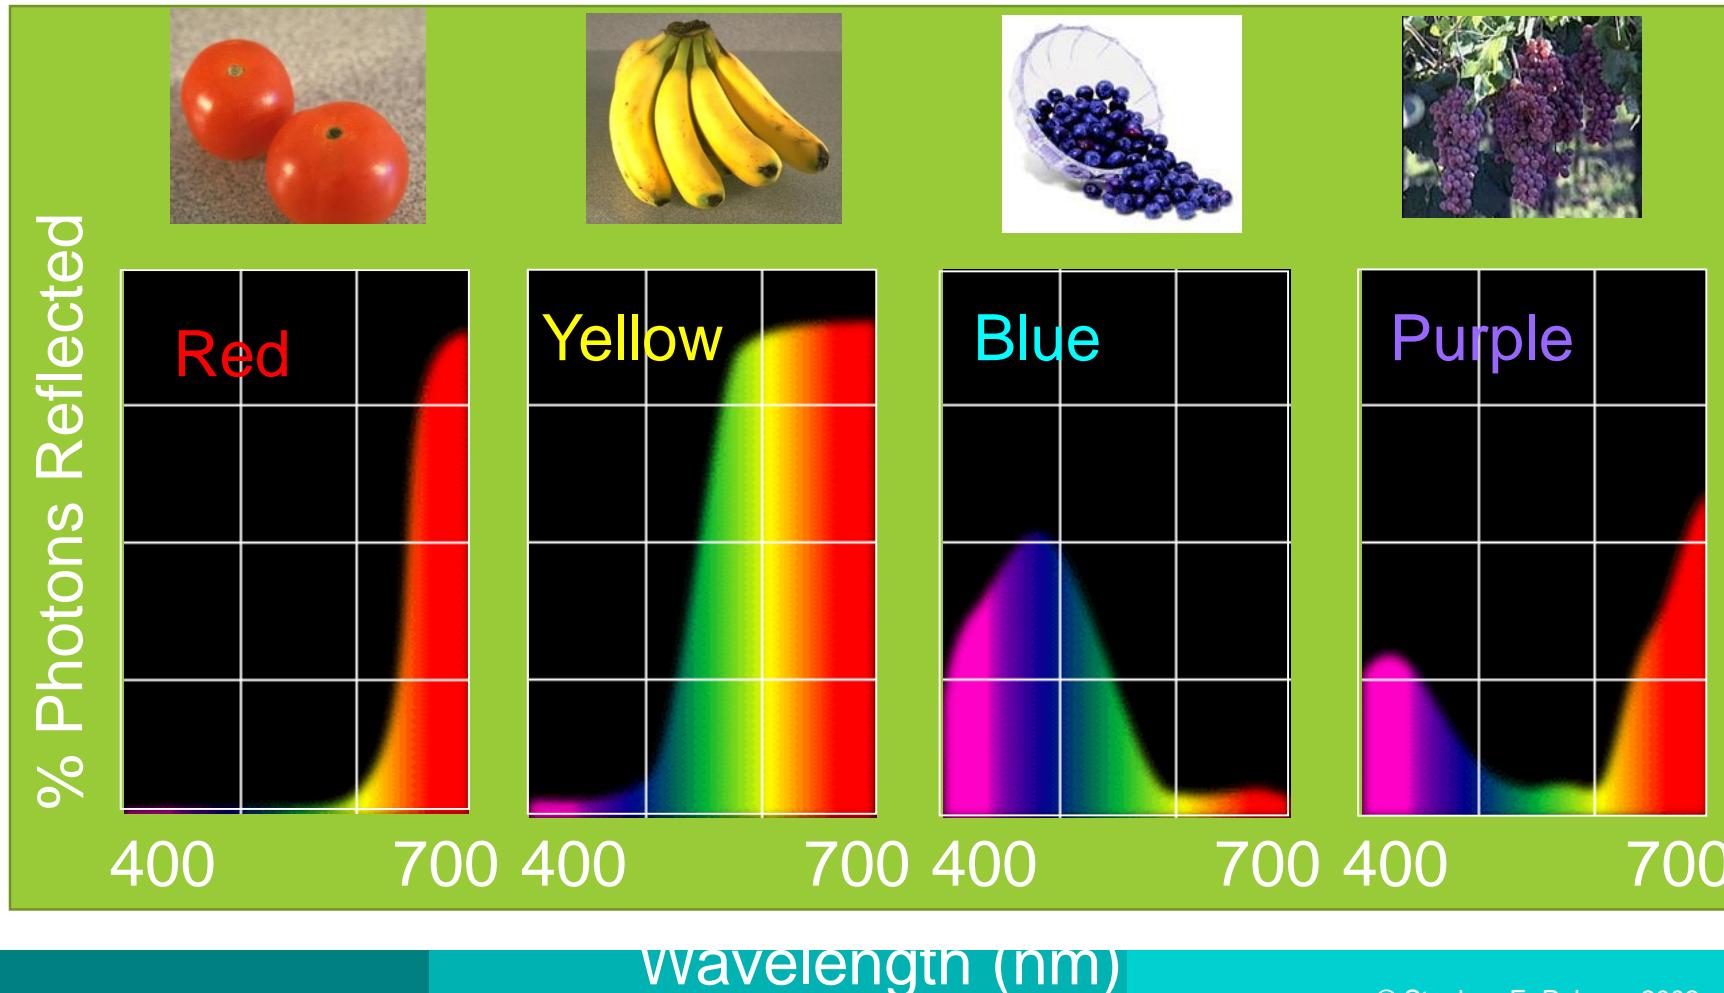
Human Luminance Sensitivity Function

Any patch of light can be completely described physically by its spectrum: the number of photons (per time unit) at each wavelength 400 - 700 nm.



The Physics of Light

Some examples of the reflectance spectra of surfaces





Color Image

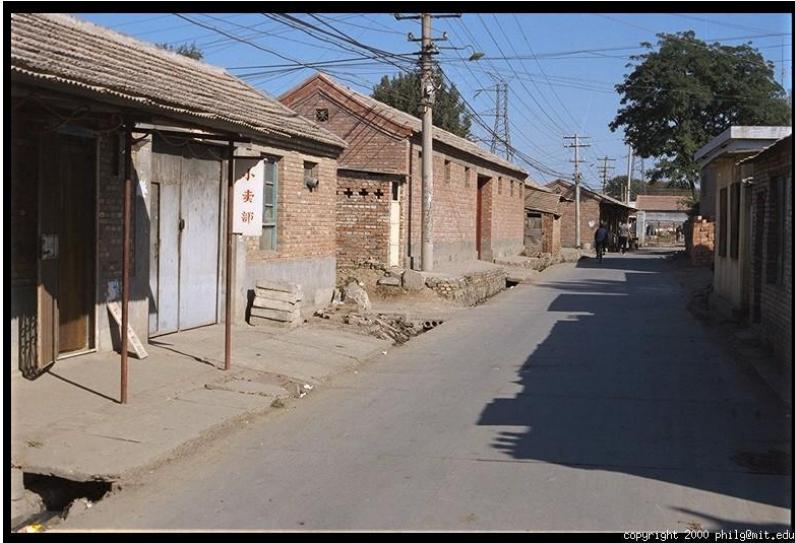
R



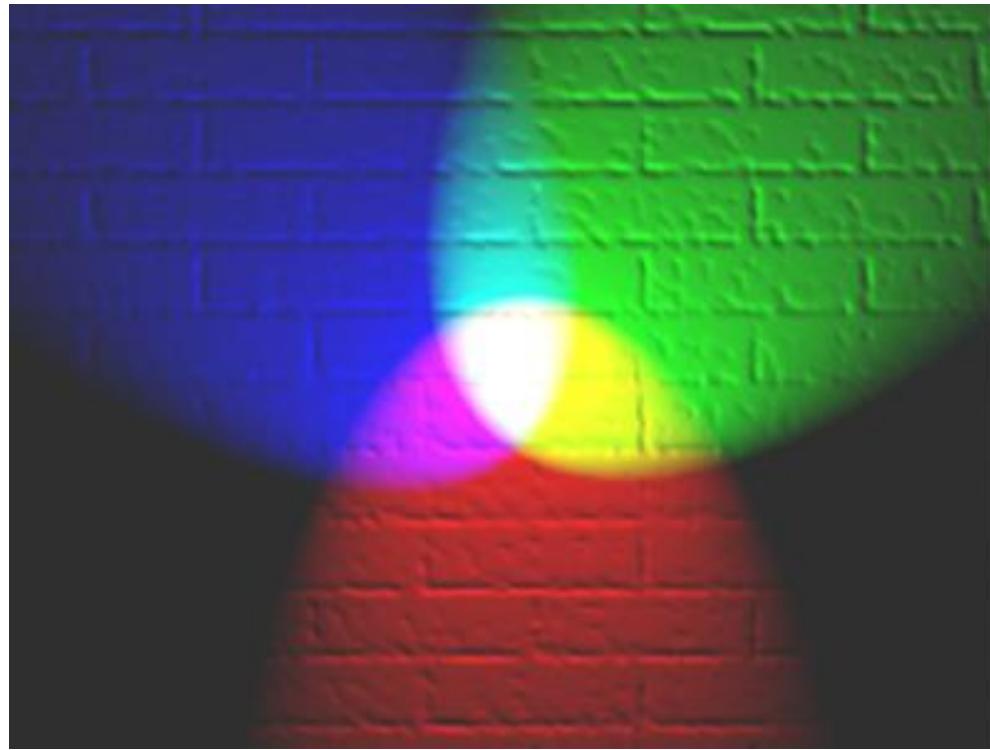
G



B



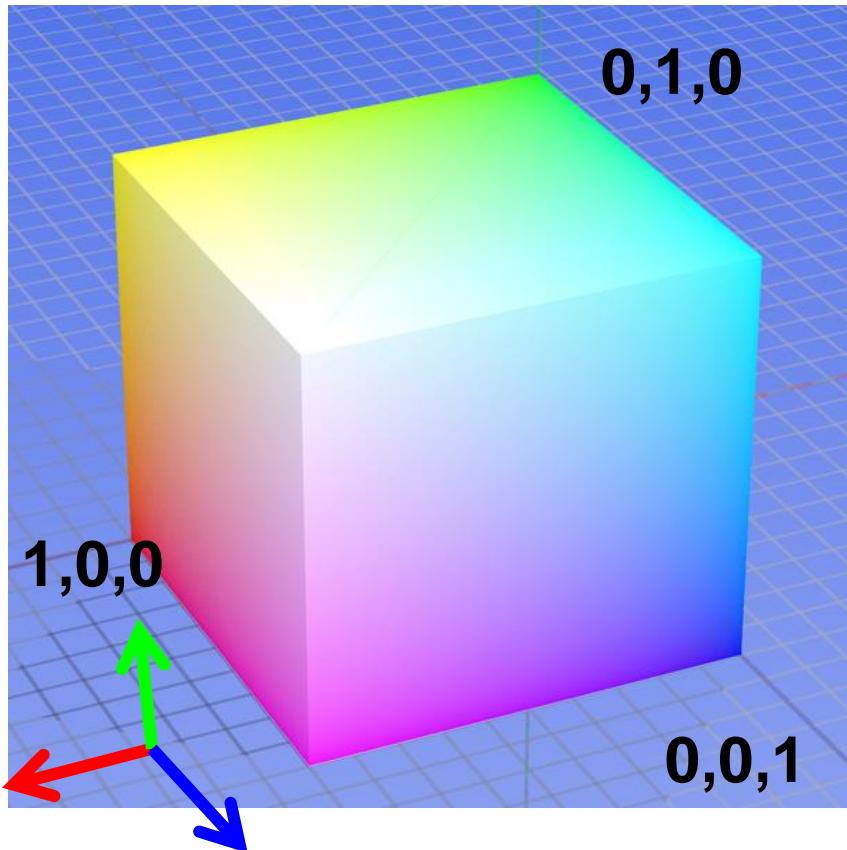
- How can we represent color?





Color spaces: RGB

Default color space



Some drawbacks

- Strongly correlated channels
- Non-perceptual



R
(G=0,B=0)



G
(R=0,B=0)



B
(R=0,G=0)

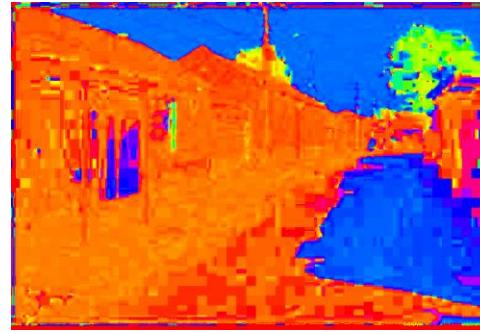
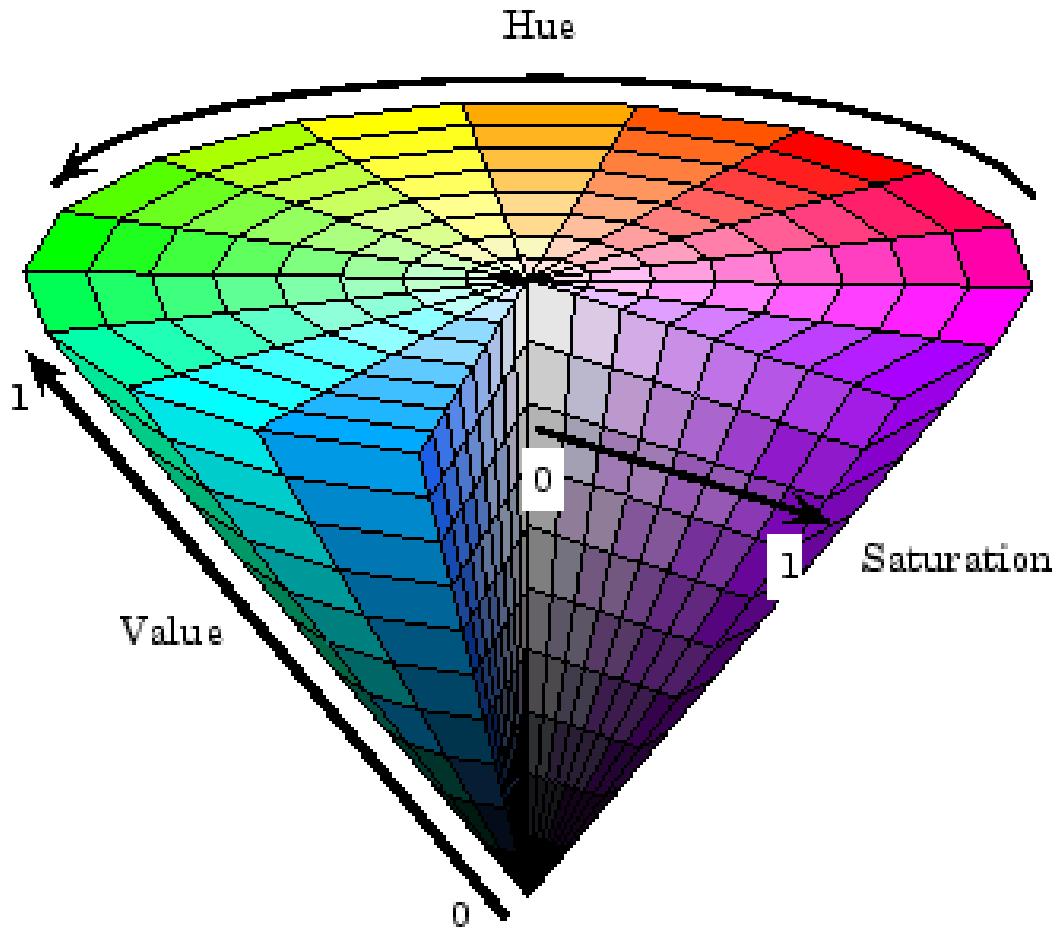




Color spaces: HSV



Intuitive color space



H
($S=1, V=1$)



S
($H=1, V=1$)

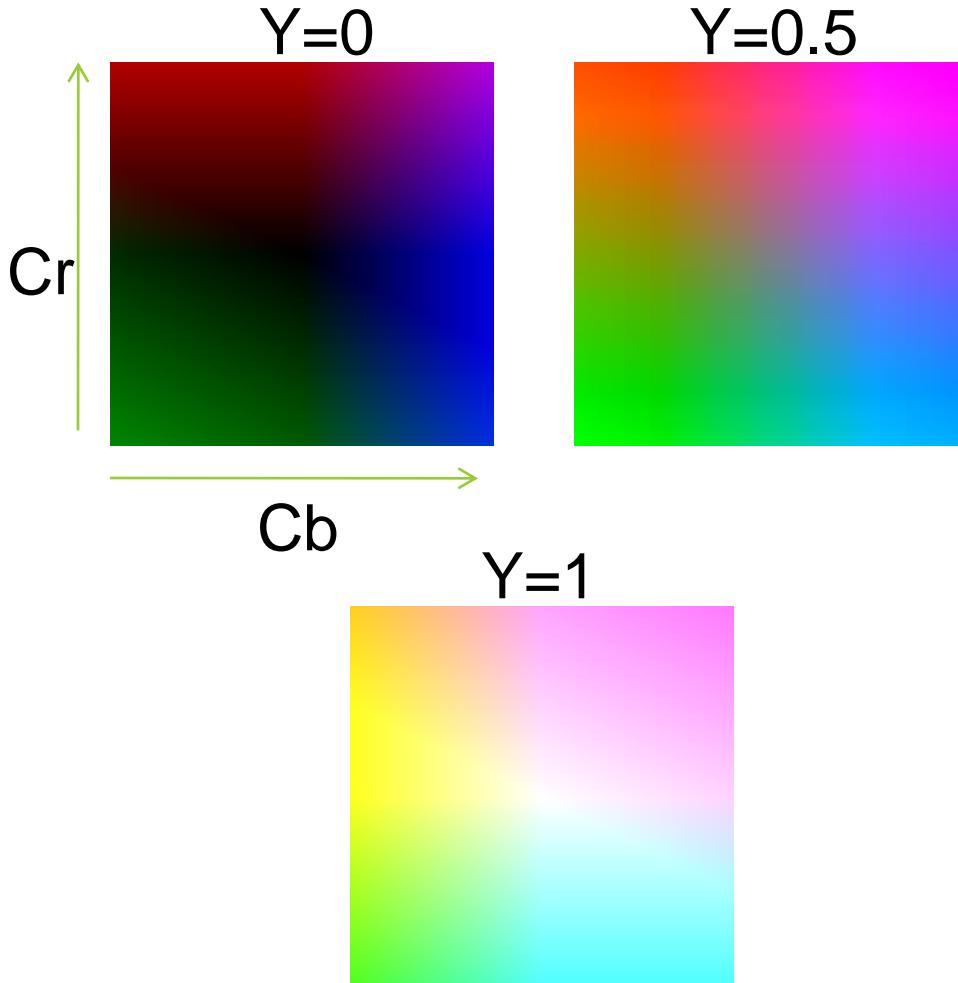


V
($H=1, S=0$)



Color spaces: YCbCr

Fast to compute, good for compression, used by TV



Y
($Cb=0.5, Cr=0.5$)



Cb
($Y=0.5, Cr=0.5$)



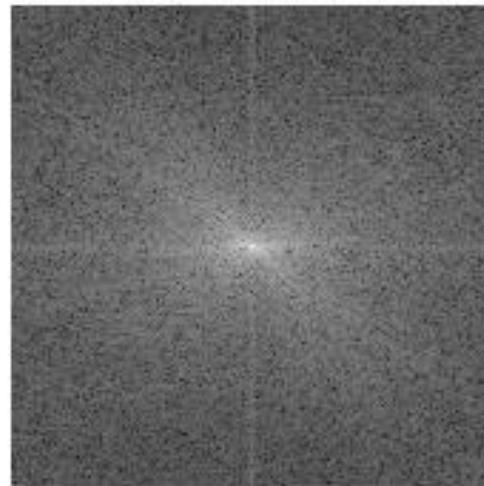
Cr
($Y=0.5, Cb=0.5$)

Any univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.

Original image



$\log(\text{FFT})$

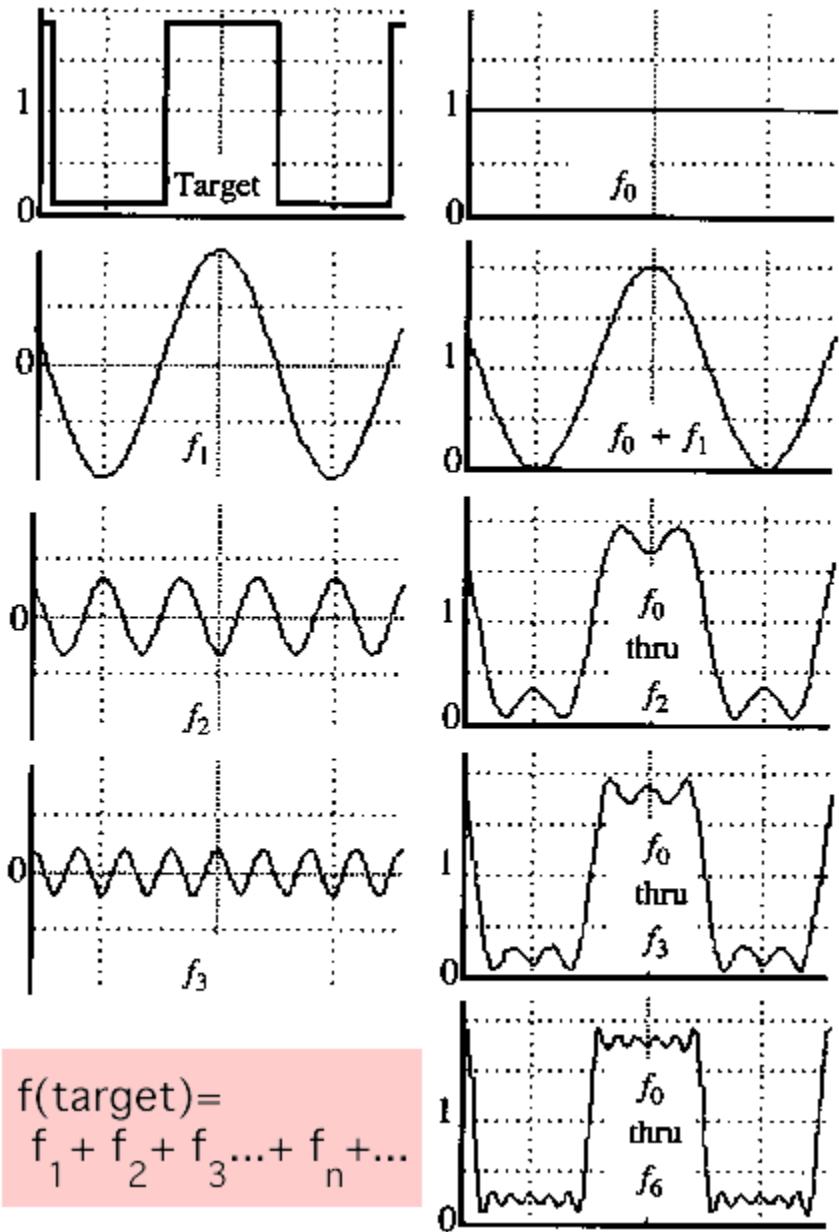


A sum of sines

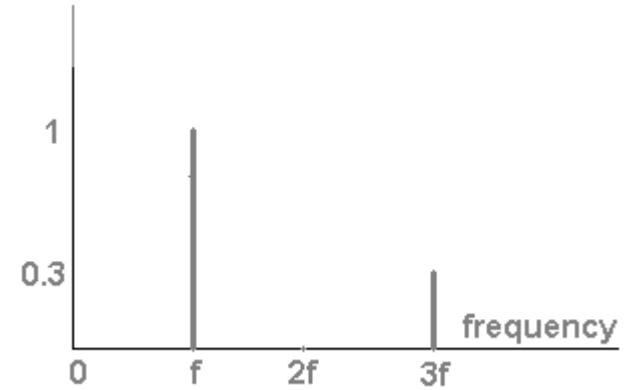
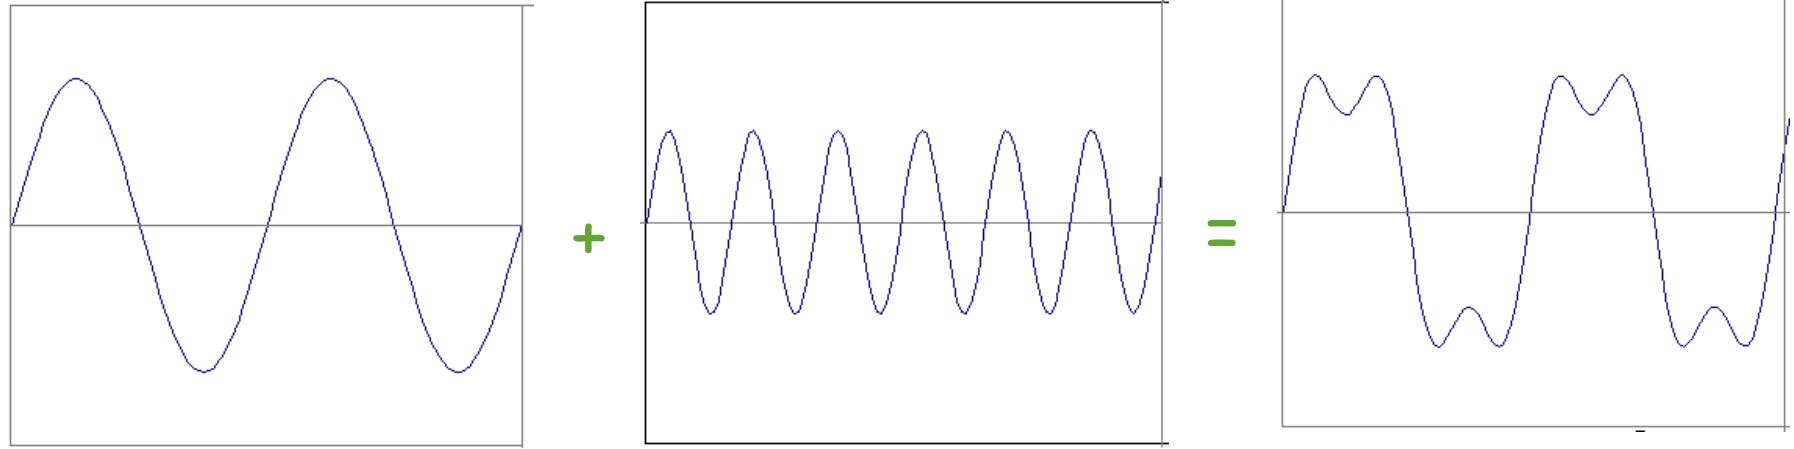
Our building block:

$$A \sin(\omega x + \phi)$$

Add enough of them to get any signal $g(x)$ you want!



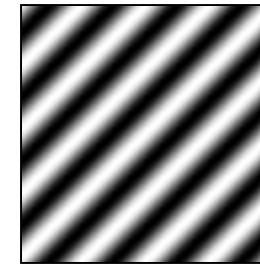
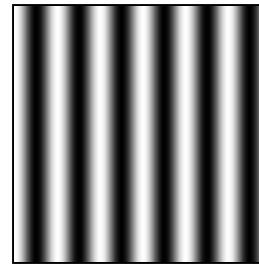
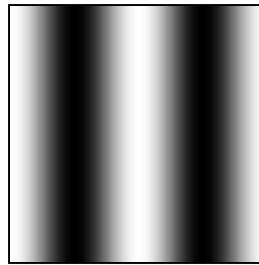
- example : $g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f)t)$



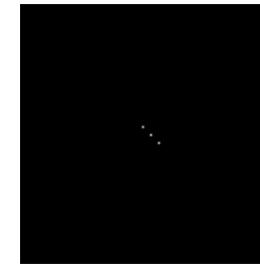
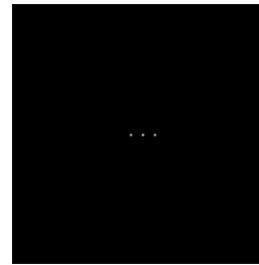
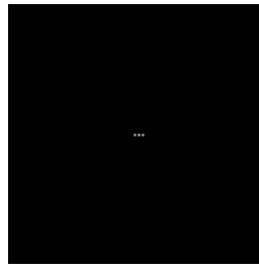


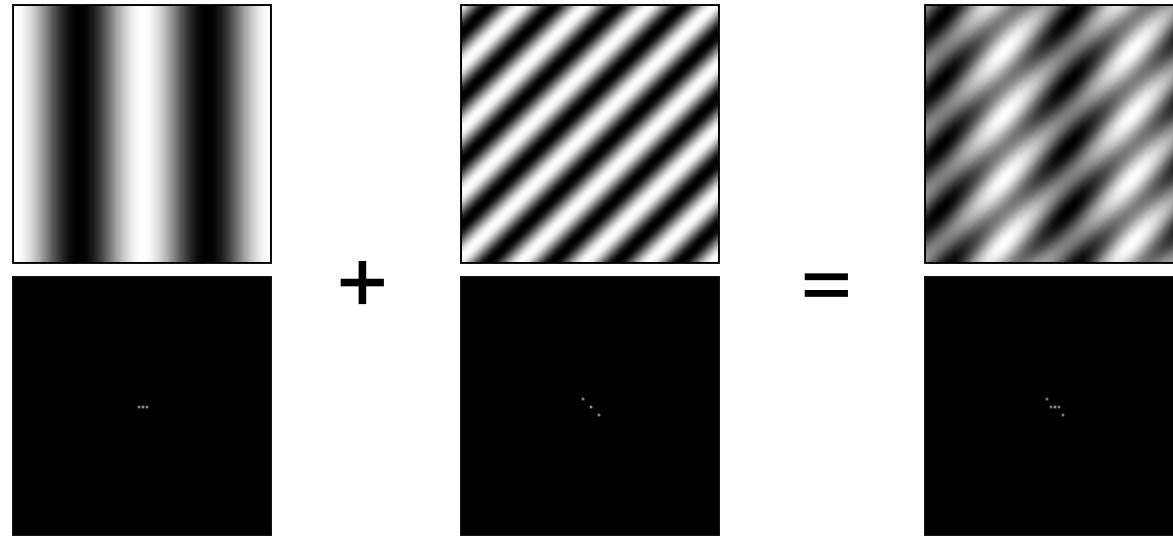
Fourier analysis in images

Intensity Image



Fourier Image

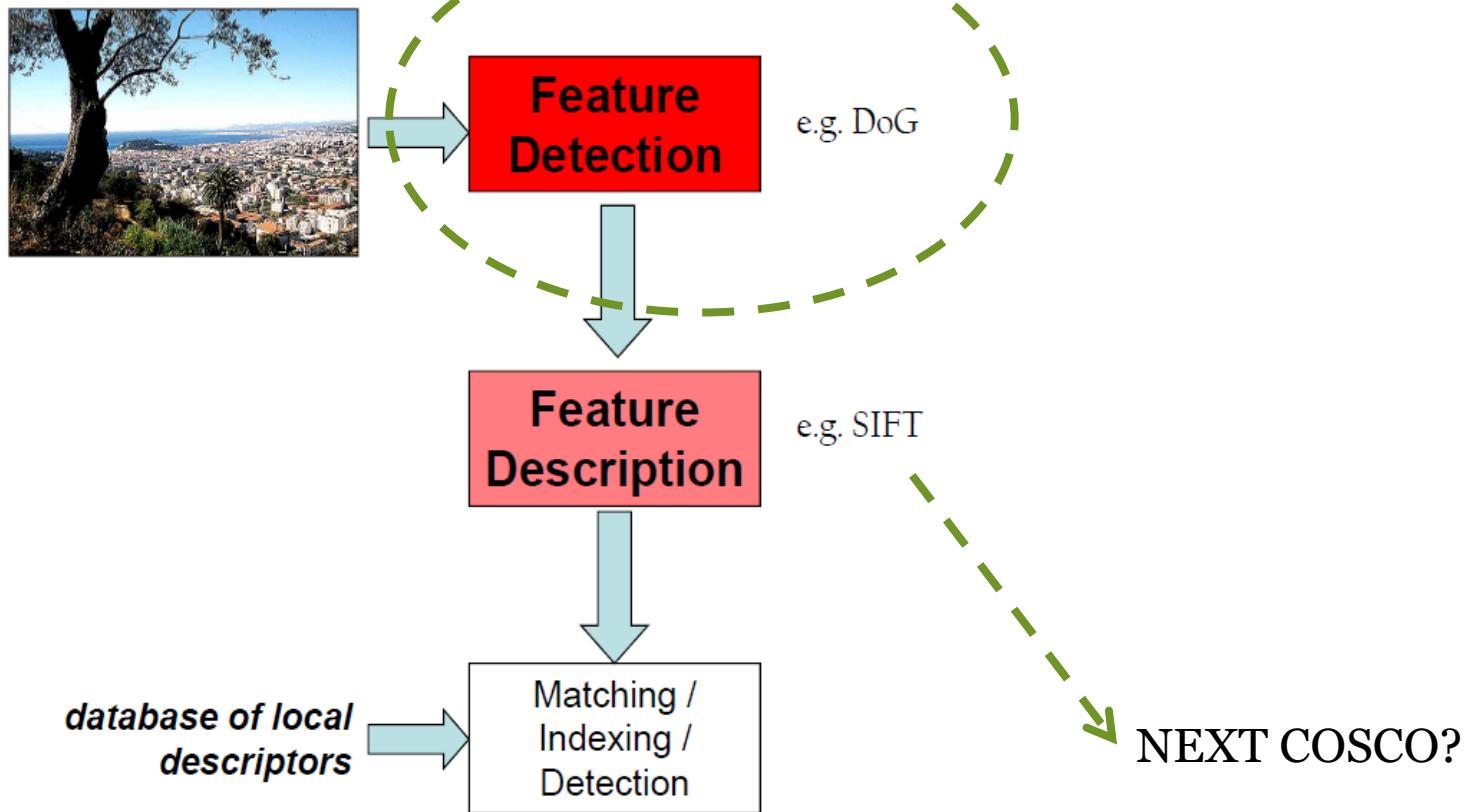




- Introduction
 - What is Vision Processing?
- Basic Knowledge
 - Image formation
 - Transformation
- Low level Algorithms
 - Filtering
 - Border Detection
 - Edge Detection
- Conclusion

- Filtering
- Simple Block matching.
- Border Detection
- Edge Detection

The big picture...



Slide credit Fei Fei Li



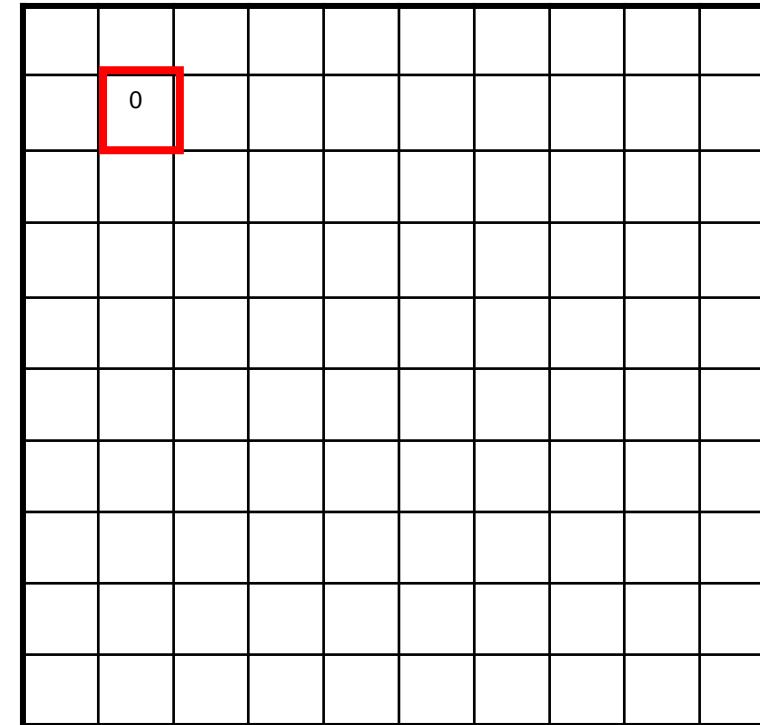
Image Filtering

- Compute function of local neighborhood at each position.
- Why:
 - Enhance images:
 - Denoise, resize, increase contrast, so on.
 - Extract information from images
 - Texture, edges, distinctive points, so on.
 - Detect patterns:
 - Template matching

$$g[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

 $f[\cdot, \cdot]$ $h[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	90	90	90	90	90	0
0	0	0	0	90	90	90	90	90	0
0	0	0	0	90	90	90	90	90	0
0	0	0	0	90	0	90	90	90	0
0	0	0	0	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0



$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

Credit: S. Seitz



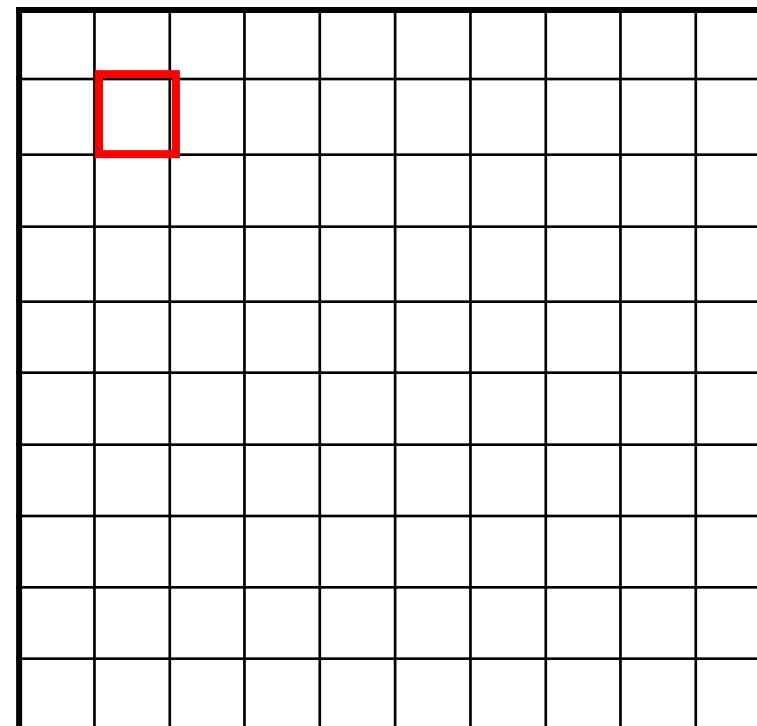
Image Filtering (3)

$f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

$$g[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

Credit: S. Seitz



Image Filtering (4)

 $f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

 $h[\cdot, \cdot]$

$$g[\cdot, \cdot] \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

Credit: S. Seitz



Image Filtering (5)

$g[\cdot, \cdot]$

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f[., .]$

$h[., .]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

Credit: S. Seitz



Original

0	0	0
0	0	1
0	0	0



Shifted left
By 1 pixel

Source: D. Lowe



Original

0	0	0
0	2	0
0	0	0

-

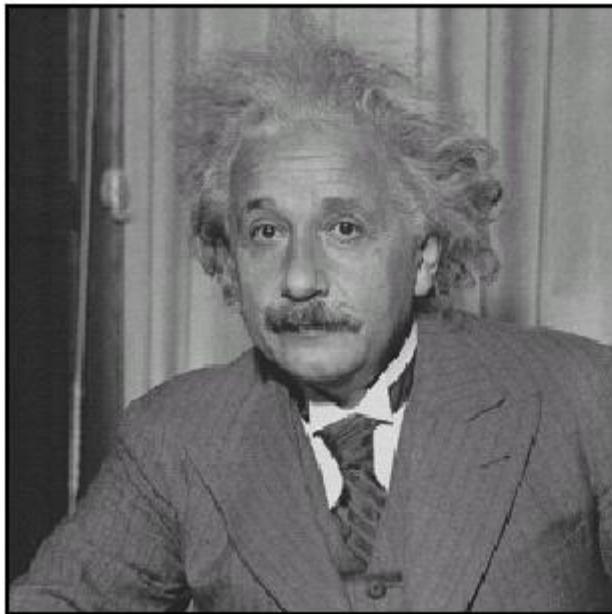
$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1



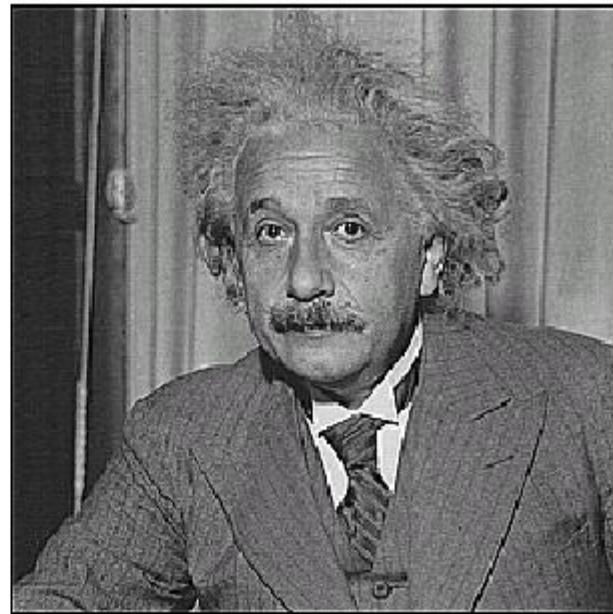
Sharpening filter

- Accentuates differences with local average

Source: D. Lowe

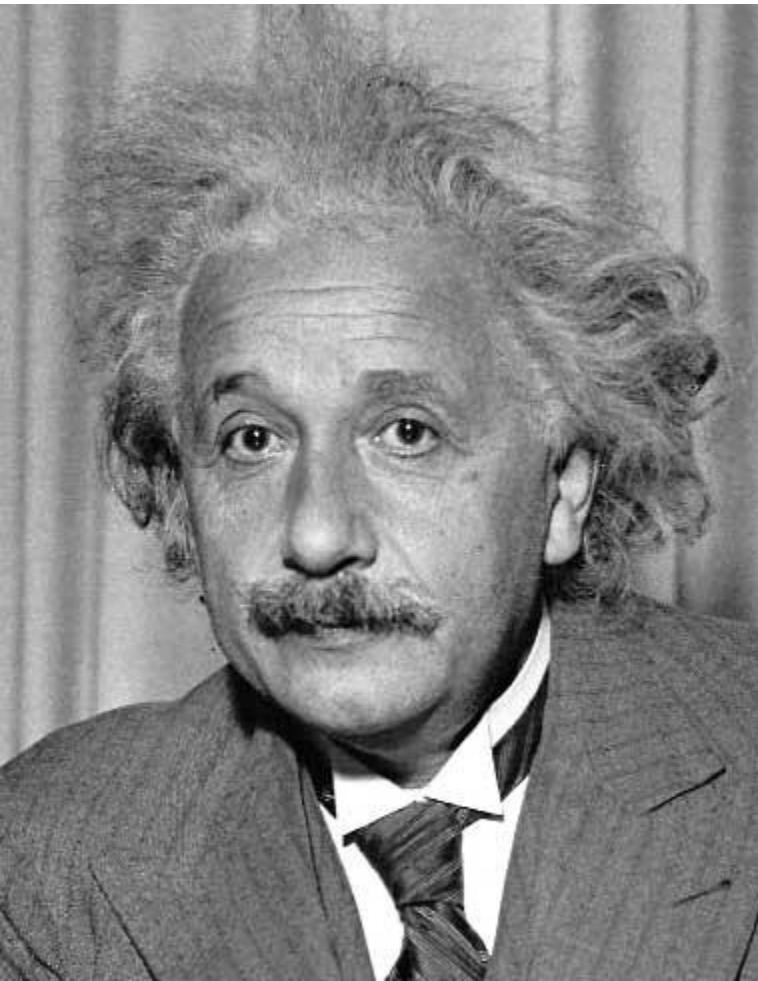


before

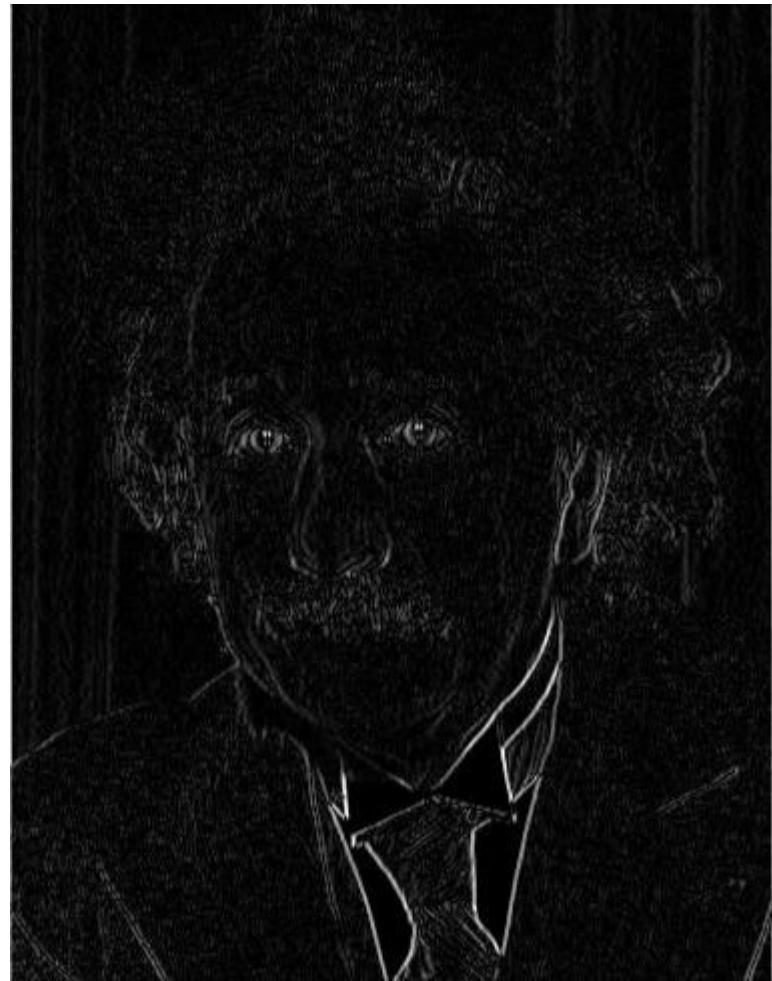


after

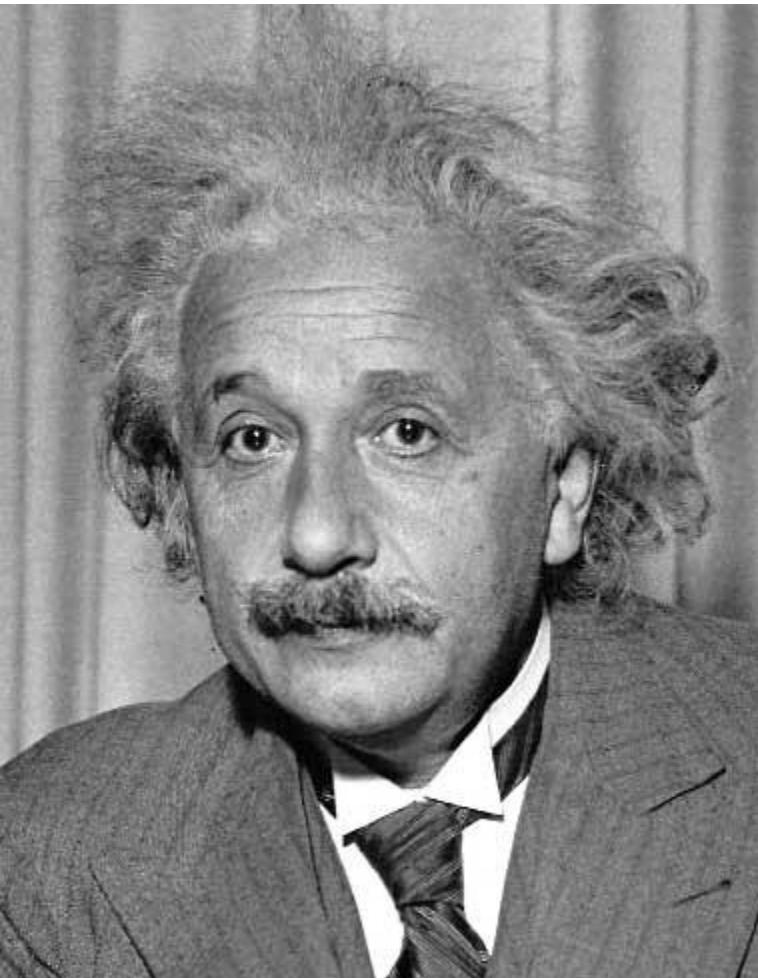
Source: D. Lowe



1	0	-1
2	0	-2
1	0	-1



Vertical Edge
(absolute value)

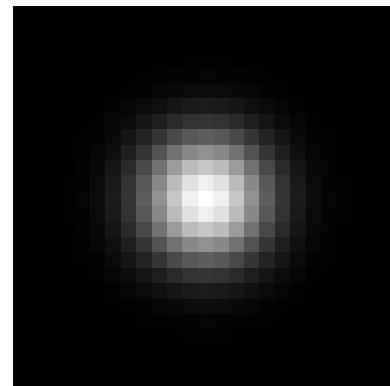
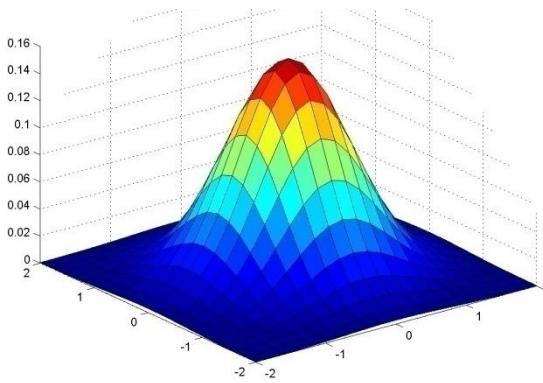


1	2	1
0	0	0
-1	-2	-1



Horizontal Edge
(absolute value)

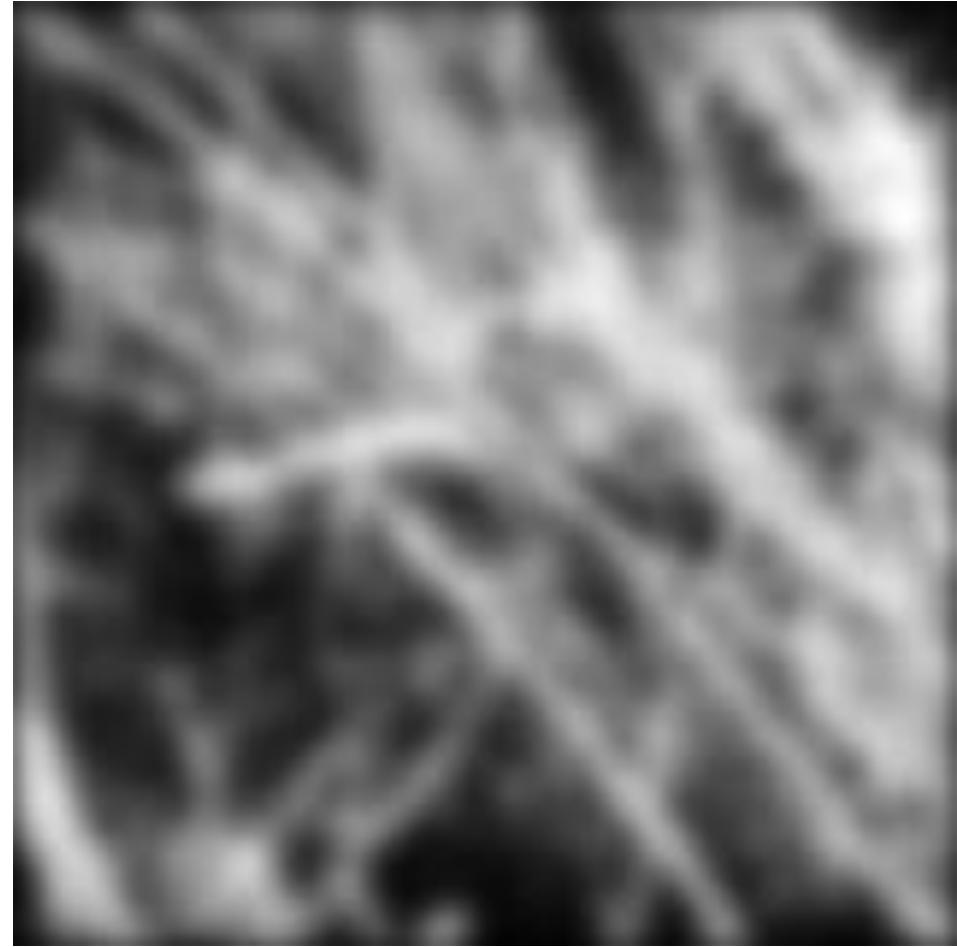
- Weight contributions of neighboring pixels by nearness.
- Gaussian Filter is low pass filter.

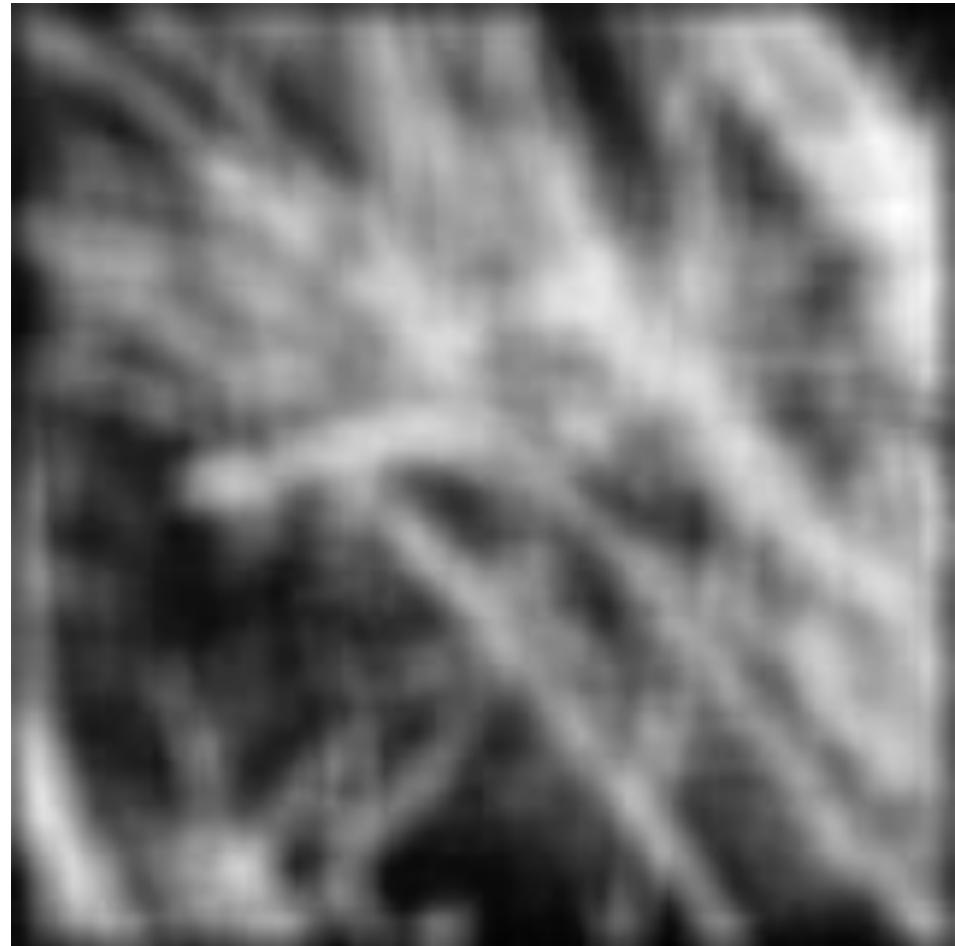


0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

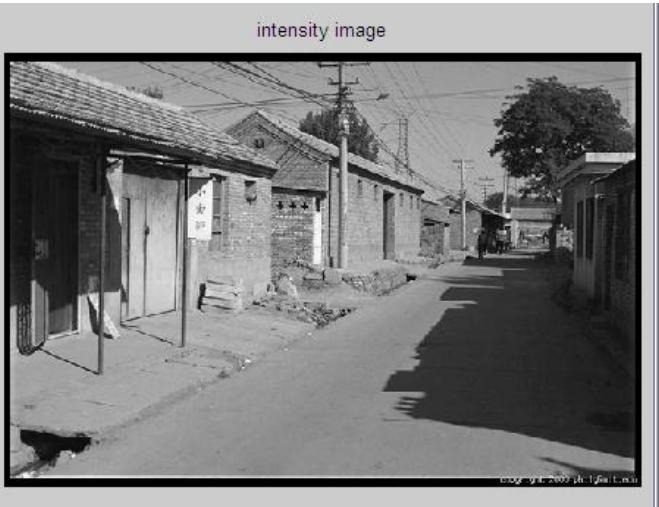
$5 \times 5, \sigma = 1$

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



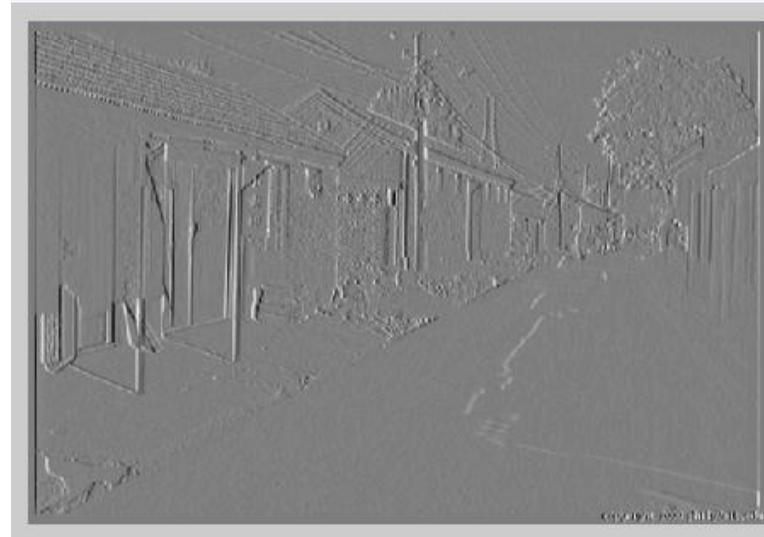


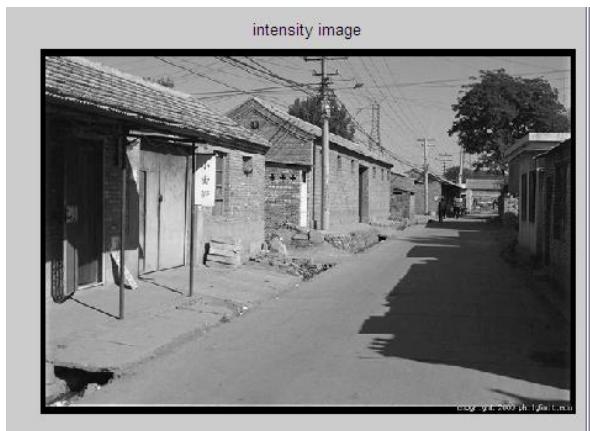
1	0	-1
2	0	-2
1	0	-1



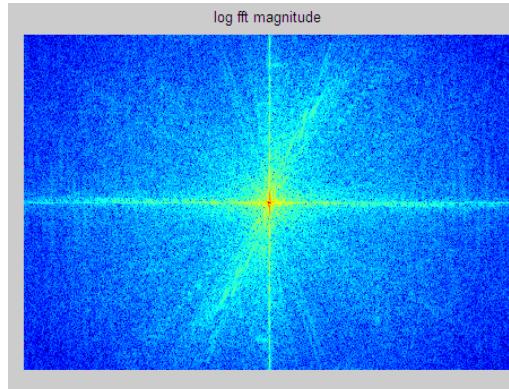
$$\ast =$$

A small grayscale image showing a vertical gradient from black to white, representing a filter kernel. It has three horizontal bars: a solid black bar, a gray bar, and a solid white bar.

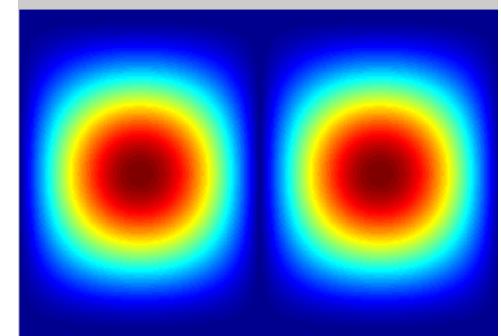




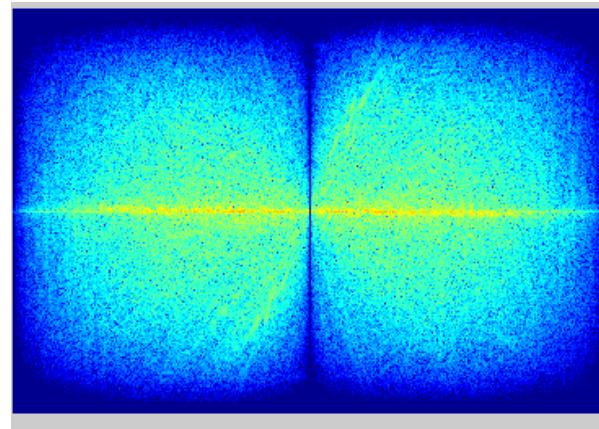
FFT



X

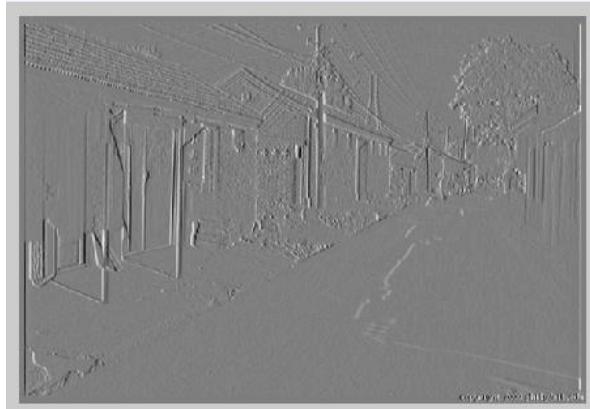


||



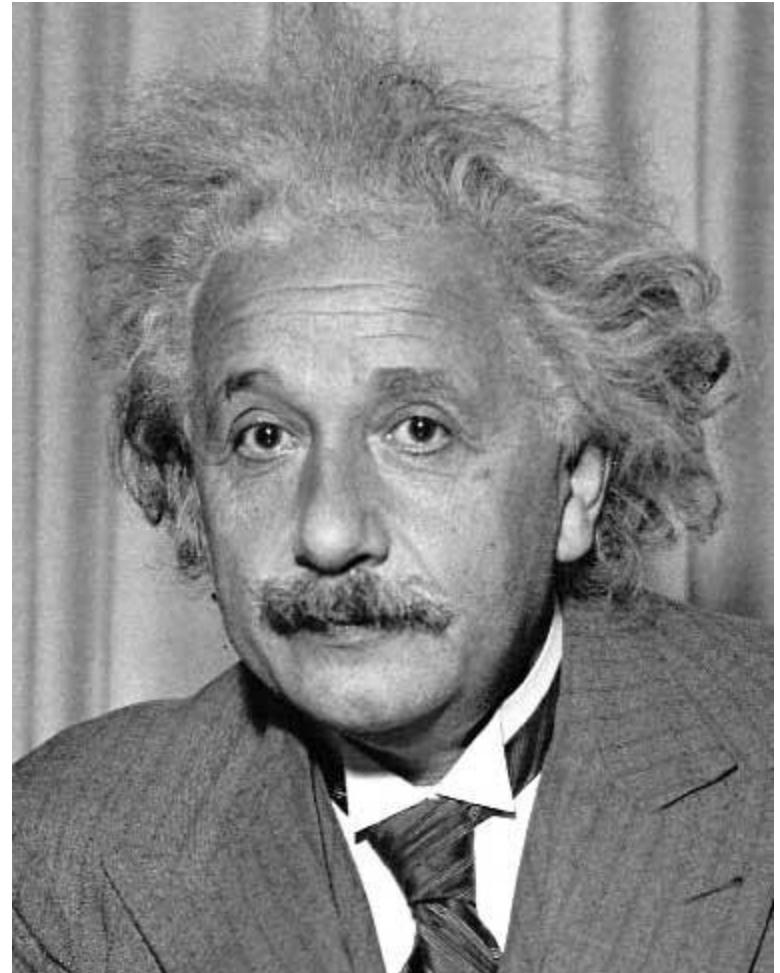
Inverse FFT

←



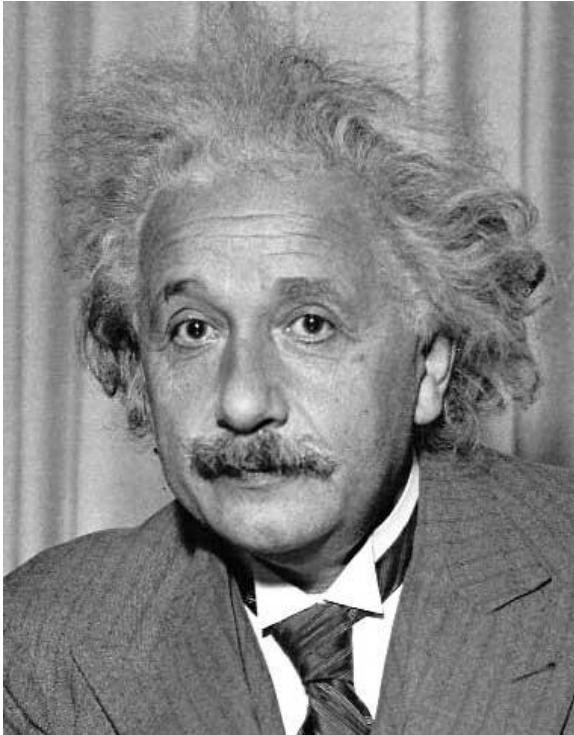
Slide: Hoiem

- Goal: find  in image
- Main challenge: What is a good similarity or distance measure between two patches?
 - Correlation
 - Zero-mean correlation
 - Sum Square Difference
 - Normalized Cross Correlation



- Goal: find  in image
- Method : SSD

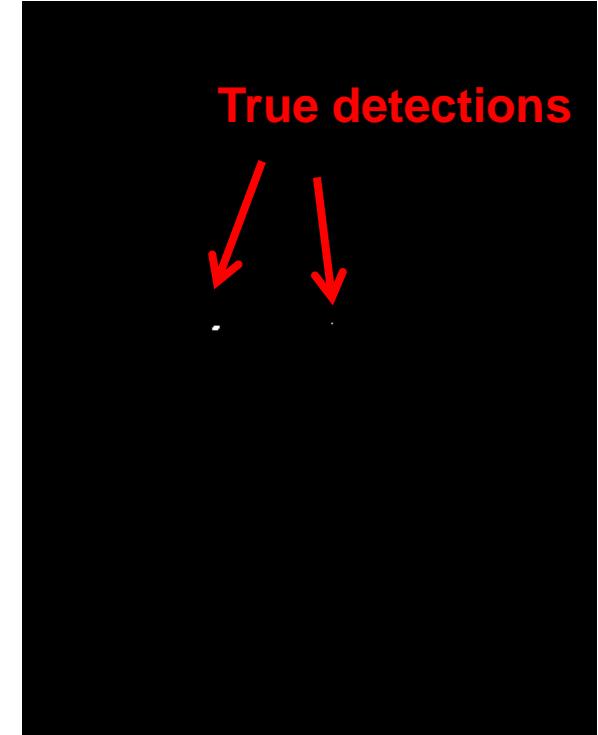
$$h[m,n] = \sum_{k,l} (g[k,l] - f[m+k, n+l])^2$$



Input



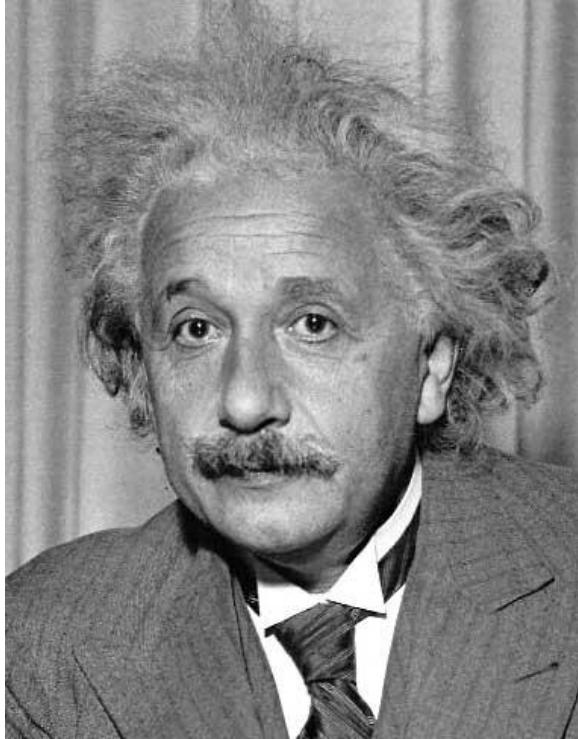
1 - $\text{sqrt}(\text{SSD})$



Thresholded Image

True detections

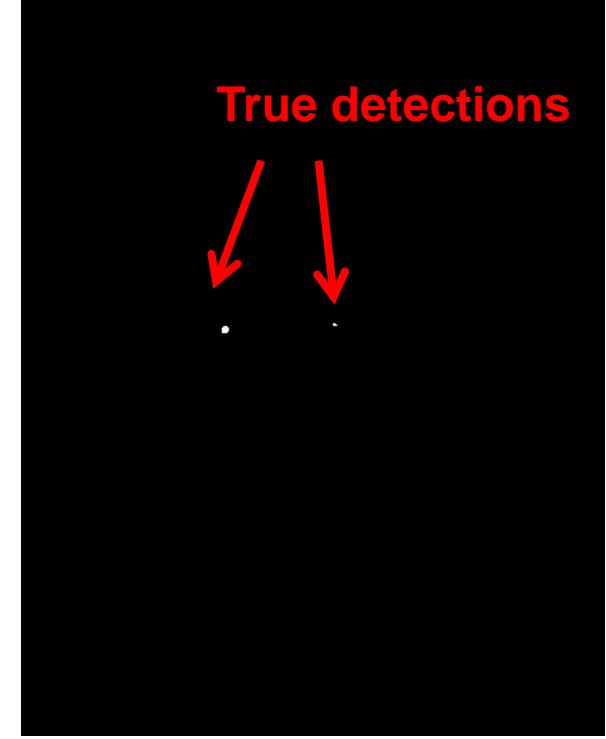
- Goal: find  in image
- Method 3: Normalized cross-correlation



Input



Normalized X-Correlation



True detections

Thresholded Image

- Goal: find  in image
- Method : Normalized cross-correlation

$$h[m,n] = \frac{\sum_{k,l} (g[k,l] - \bar{g})(f[m-k, n-l] - \bar{f}_{m,n})}{\left(\sum_{k,l} (g[k,l] - \bar{g})^2 \sum_{k,l} (f[m-k, n-l] - \bar{f}_{m,n})^2 \right)^{0.5}}$$

mean template mean image patch

Matlab: `normxcorr2(template, im)`

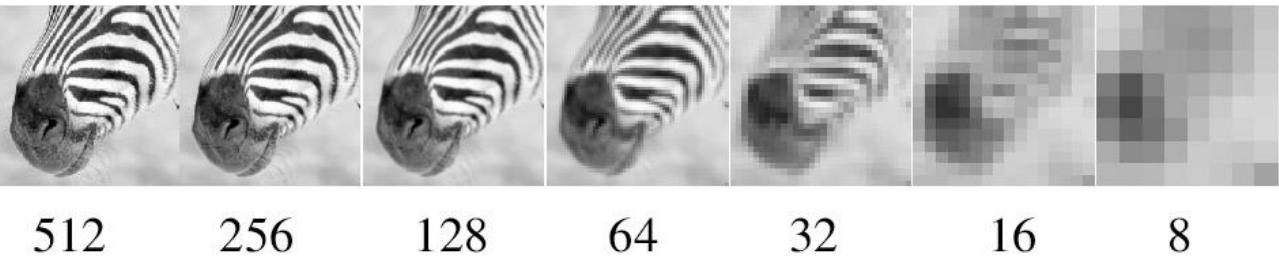
Compare

- SSD: faster, sensitive to overall intensity
- Normalized cross-correlation: slower, invariant to local average intensity and contrast

Matching smaller / larger eyes?

- Image pyramid: down-sampling and matching

Image

Low-Res
Image

Source: Forsyth

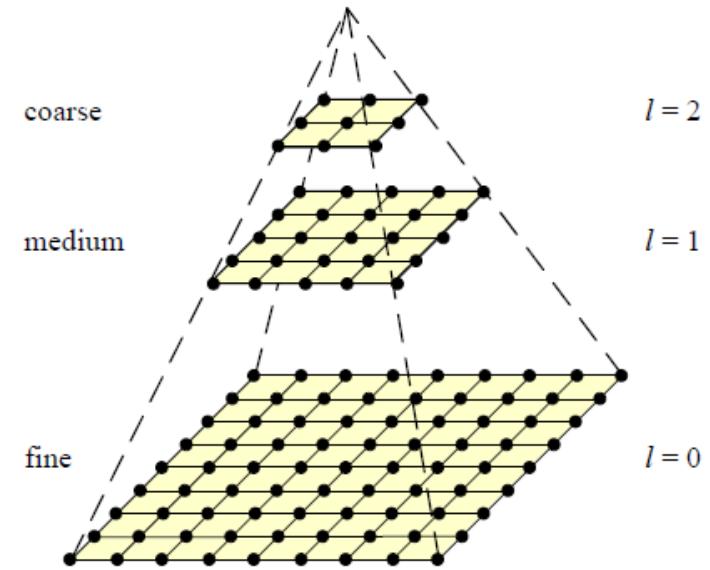


Matching with Image pyramid

Input: Image, Template

1. Match template at current scale
2. Downsample image
3. Repeat 1-2 until image is very small
4. Take responses above some threshold, perhaps with non-maxima suppression

1. Compute Gaussian pyramid
2. Align with coarse pyramid
3. Successively align with finer pyramids
 - Search smaller range

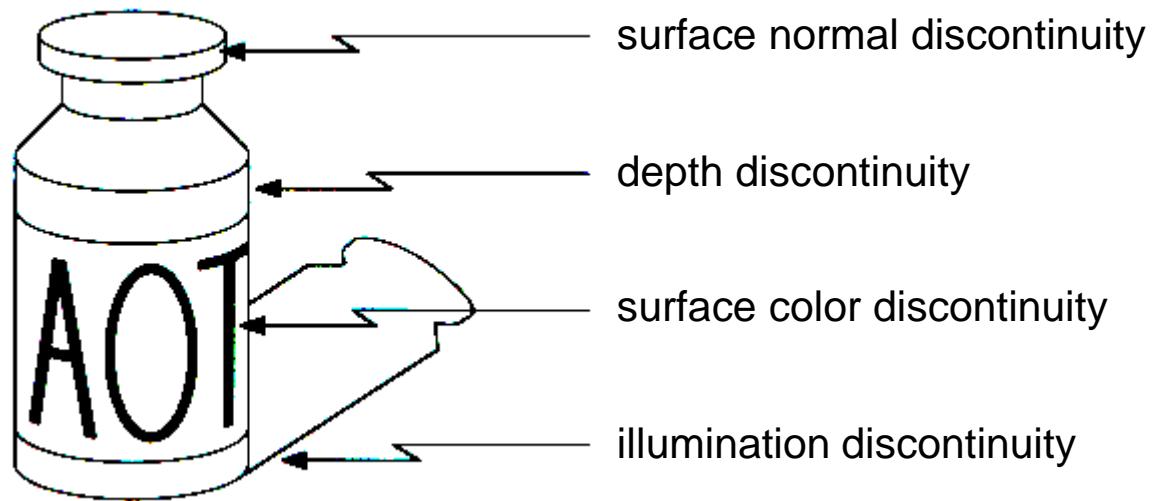


Why is this faster?

Are we guaranteed to get the same result?

- **Goal:** Identify sudden changes (discontinuities) in an image
 - Intuitively, most semantic and shape information from the image can be encoded in the edges
 - More compact than pixels
- **Ideal:** artist's line drawing (but artist is also using object-level knowledge)

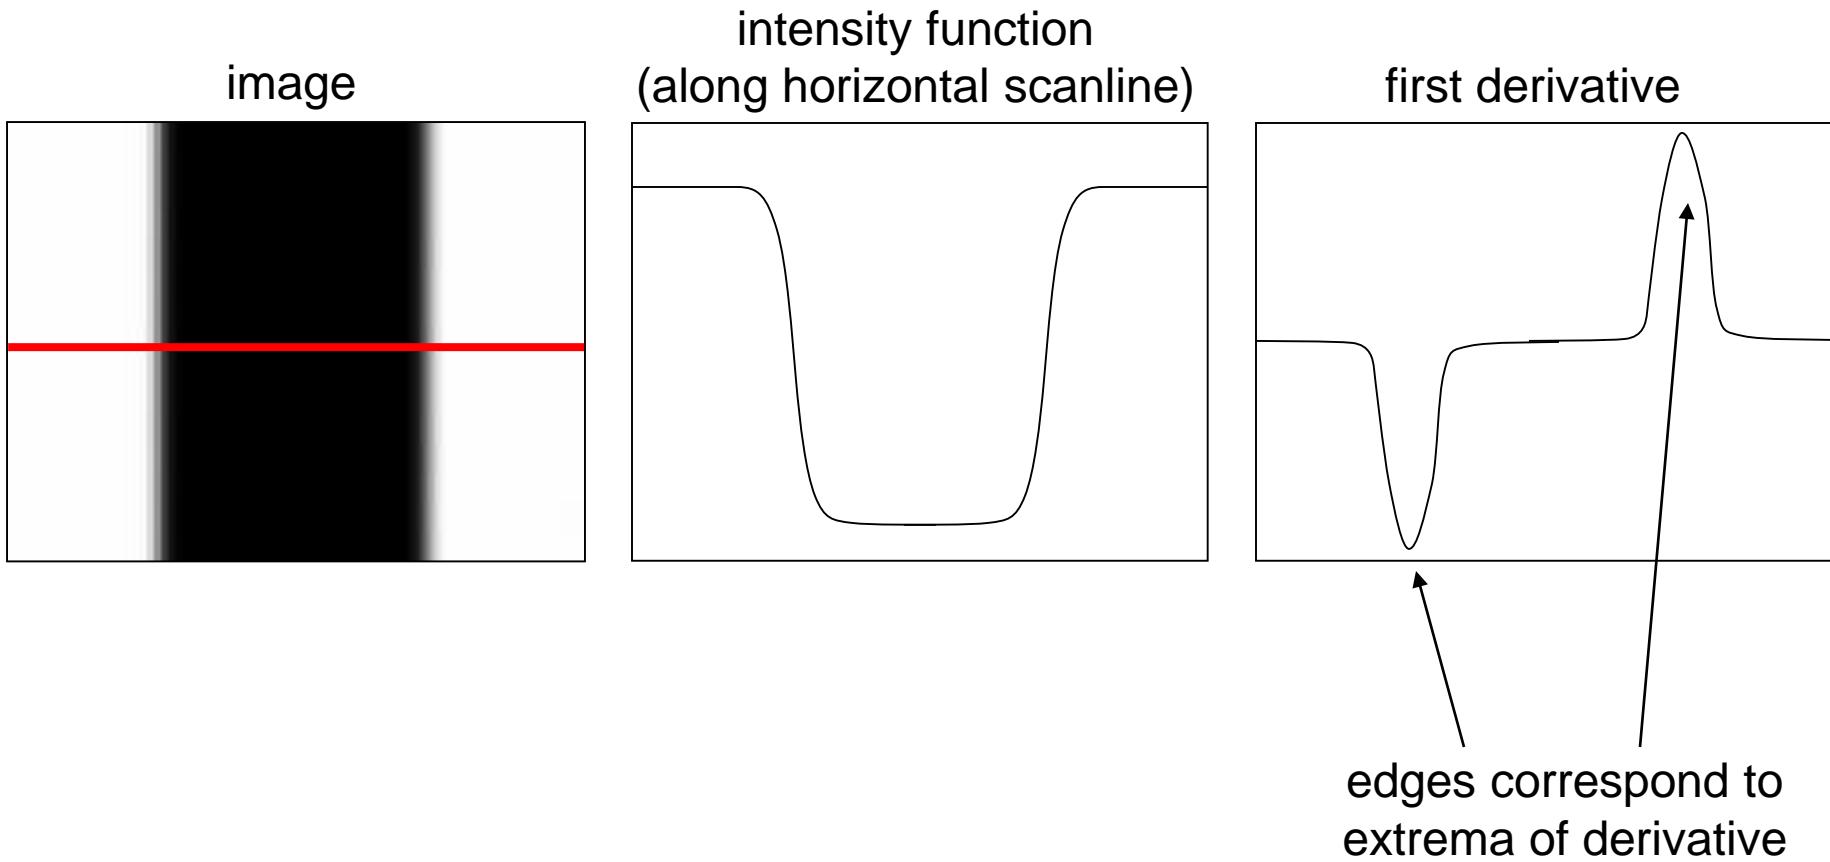




- Edges are caused by a variety of factors

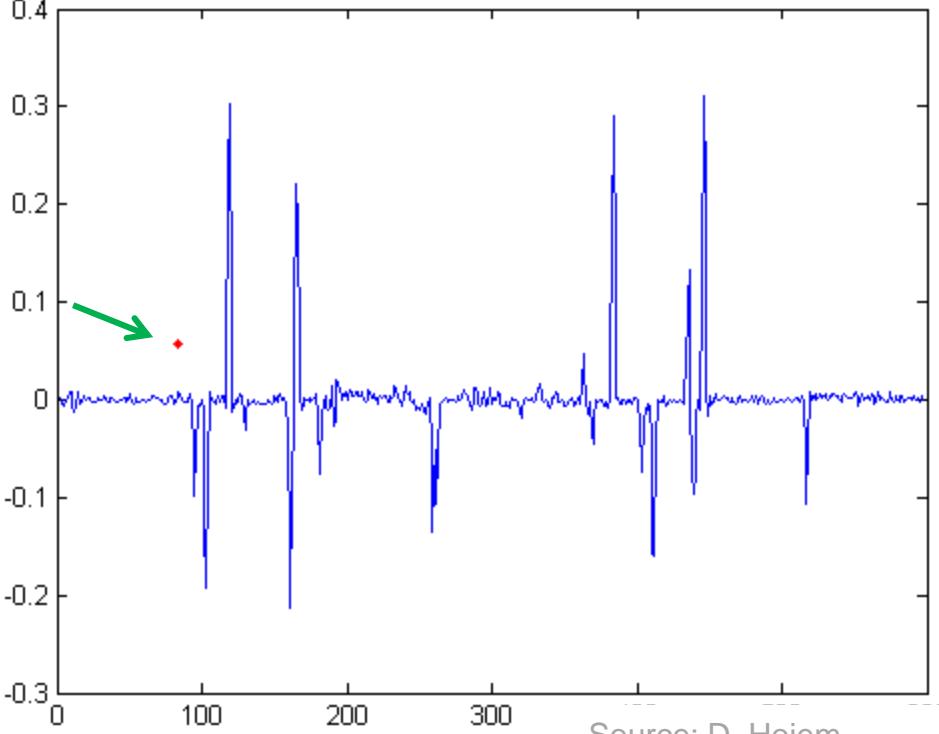
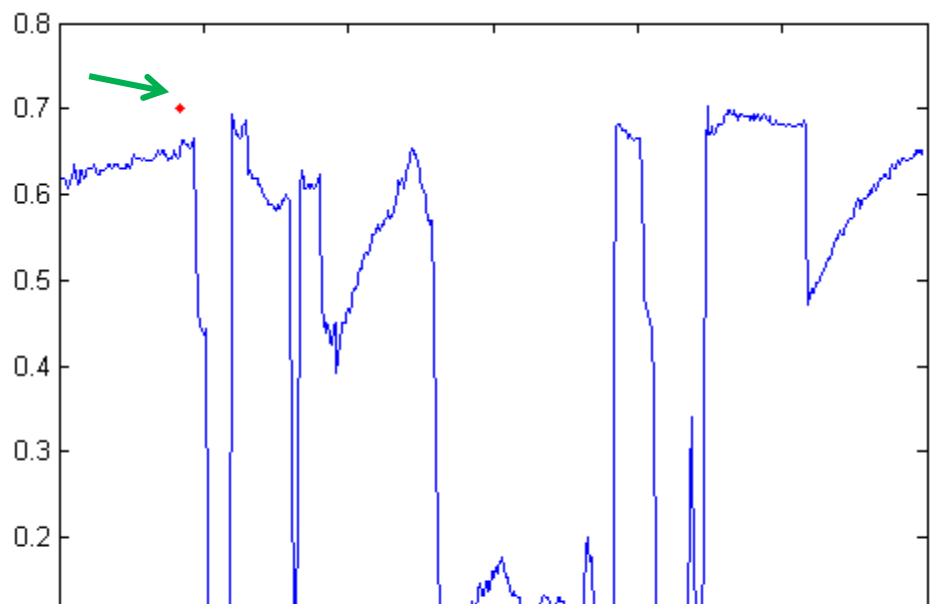
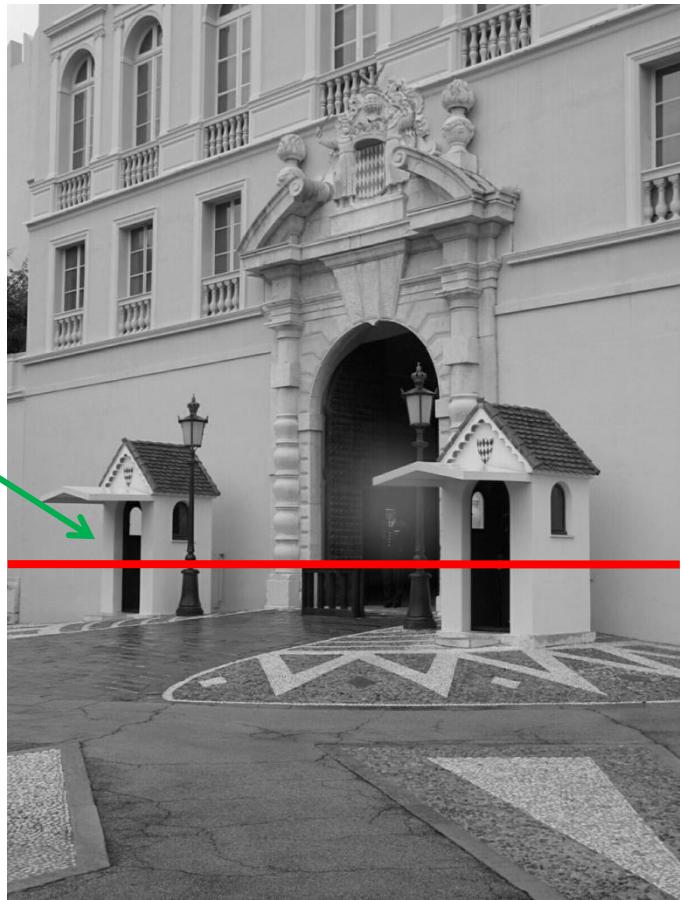
Source: Steve Seitz

- An edge is a place of rapid change in the image intensity function

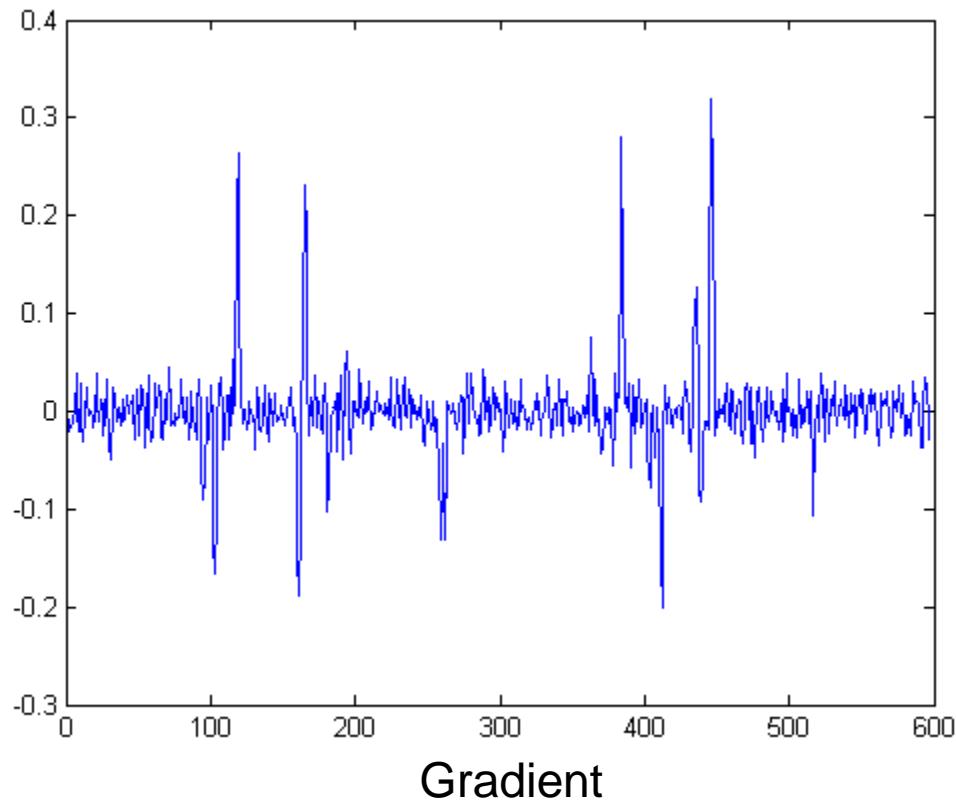
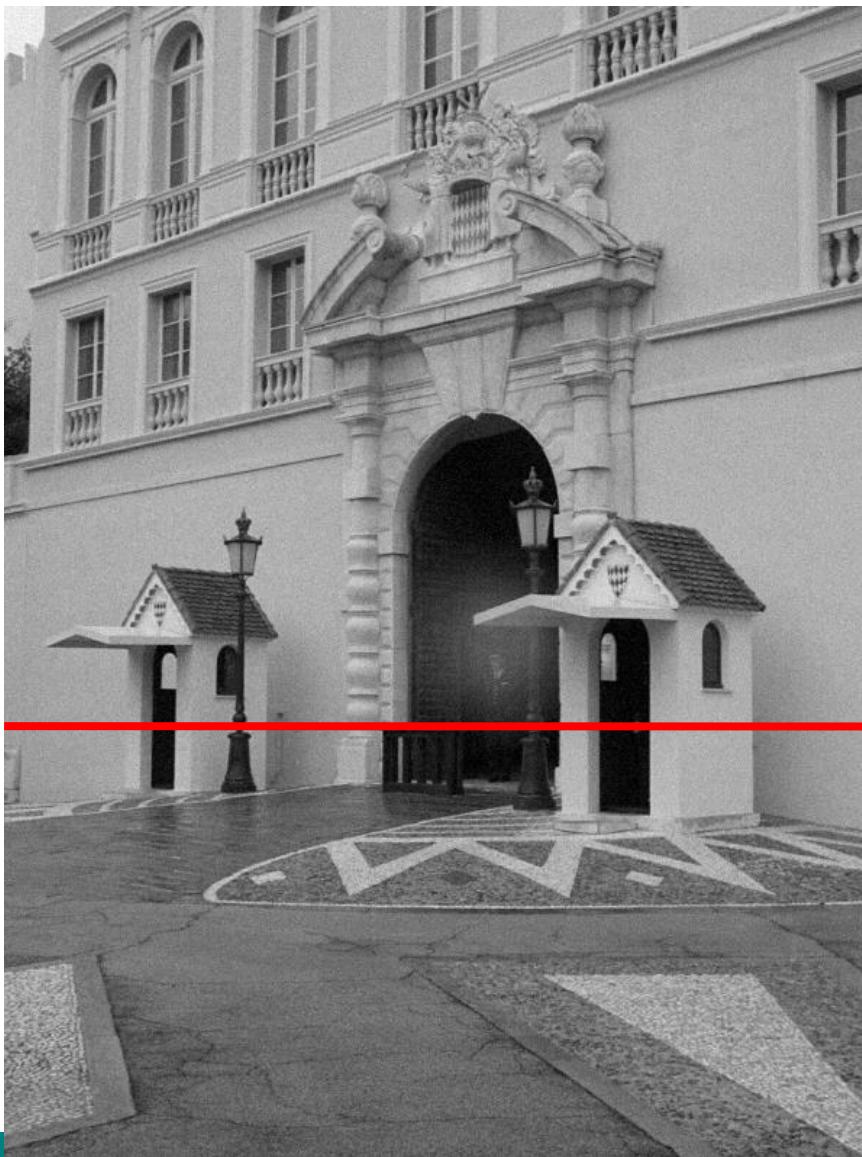




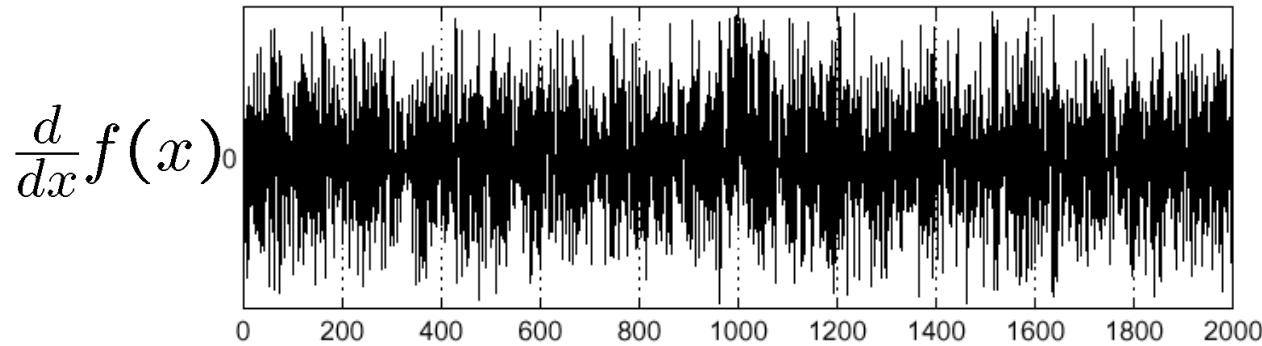
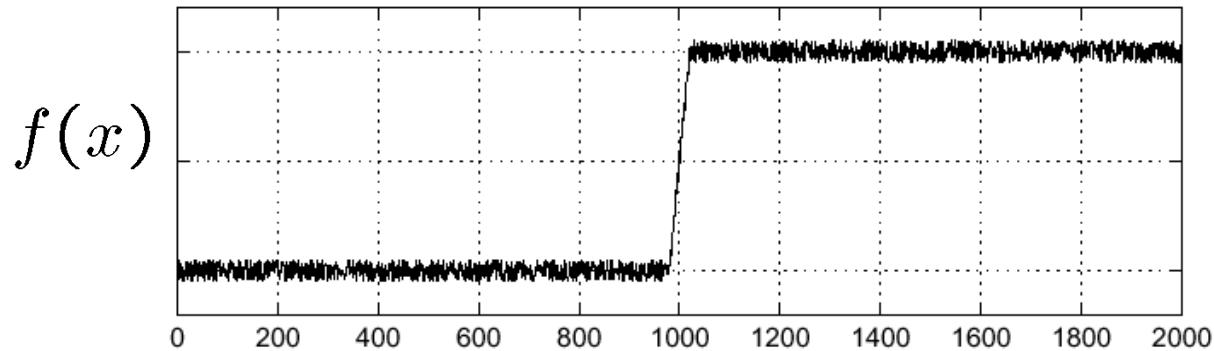
Intensity profile

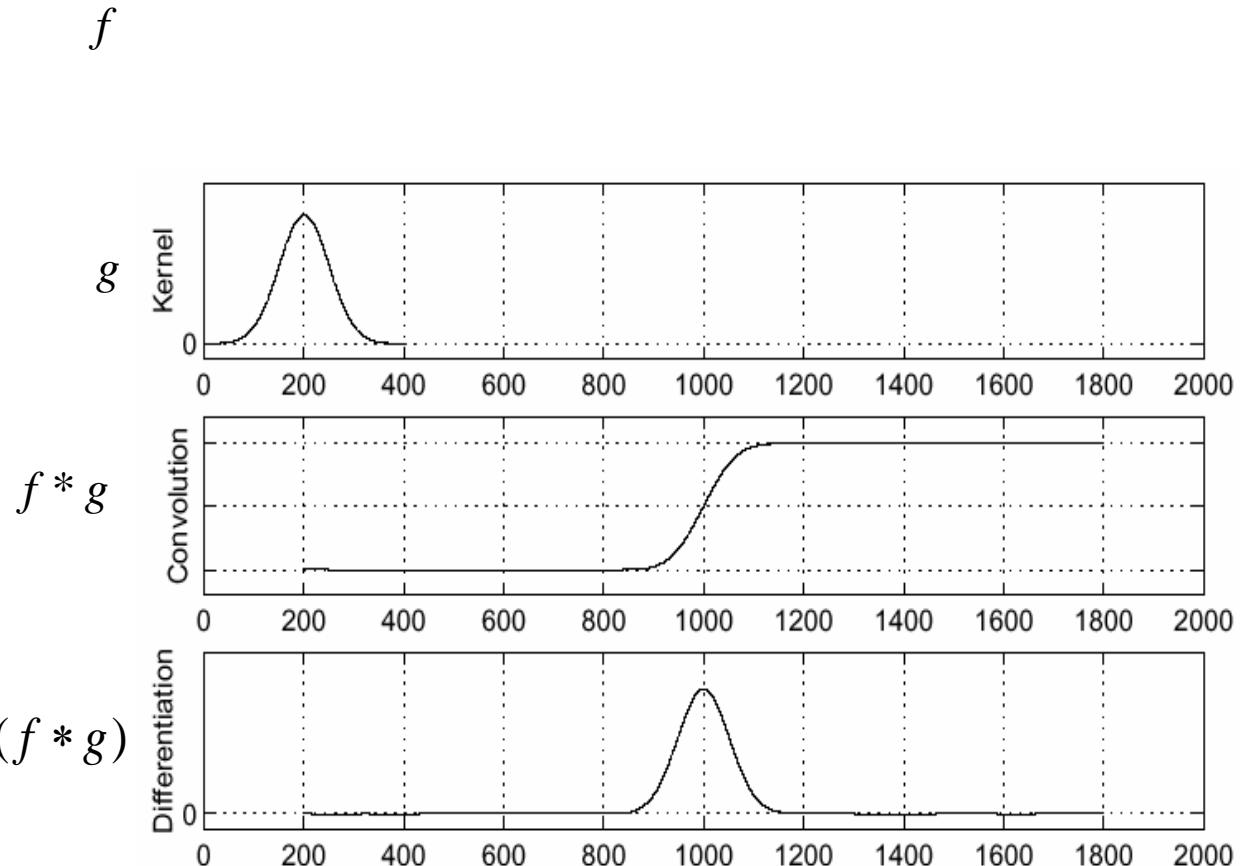


Source: D. Hoiem



- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal





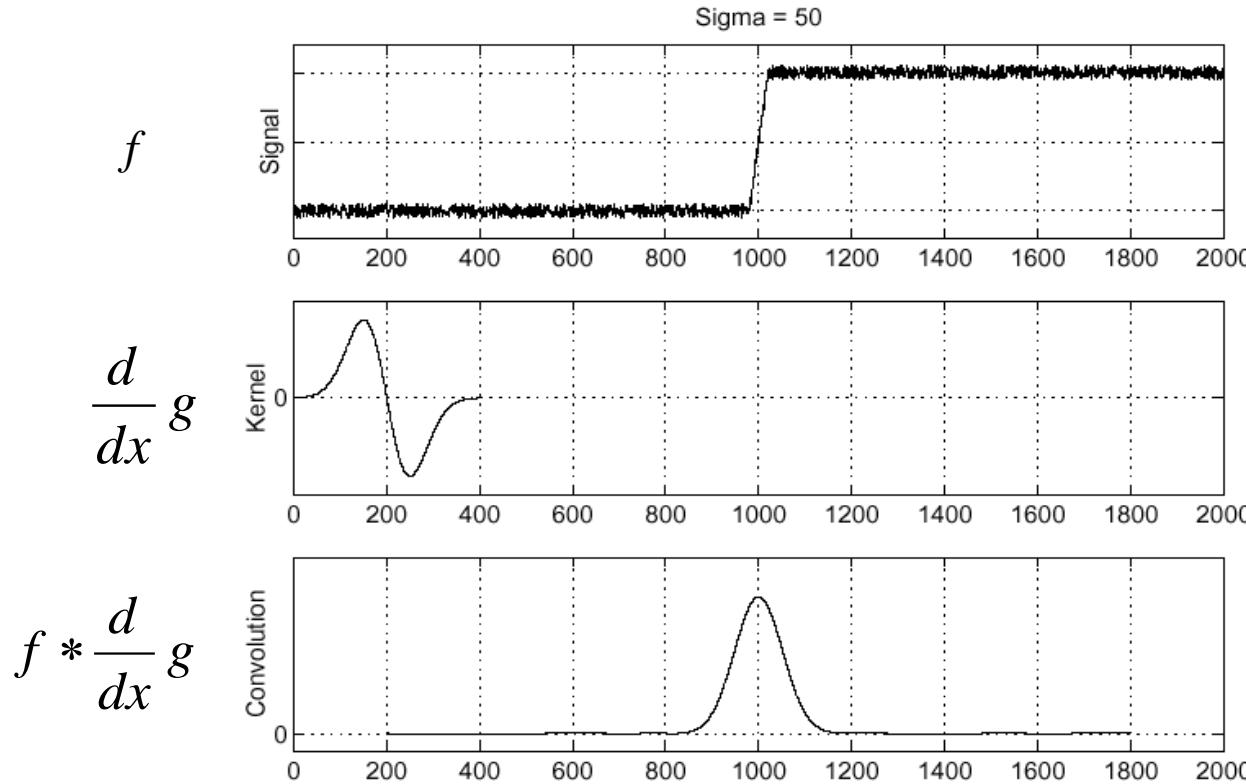
- To find edges, look for peaks in

$$\frac{d}{dx}(f * g)$$

- Differentiation is convolution, and convolution is associative:

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

- This saves us one operation:

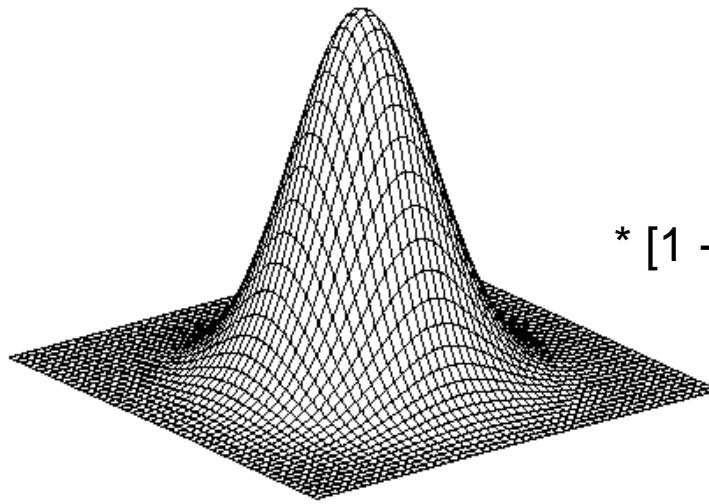




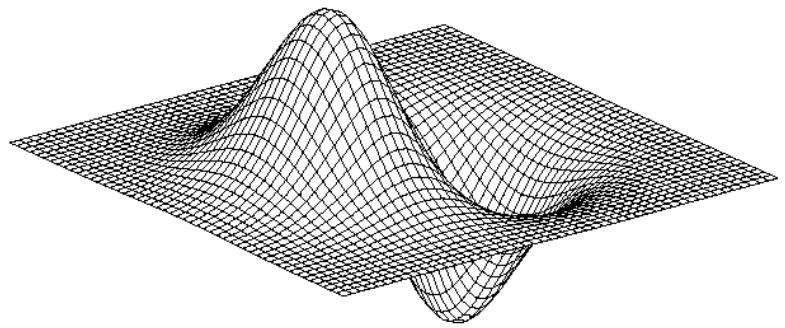
Canny edge detector

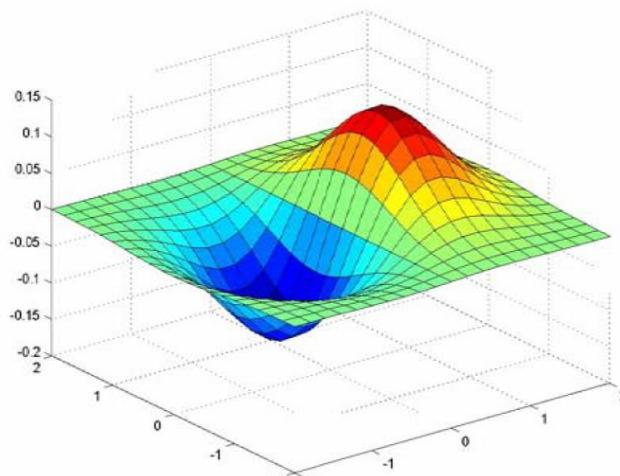
- This is probably the most widely used edge detector in computer vision
- Theoretical model: step-edges corrupted by additive Gaussian noise
- Canny has shown that the first derivative of the Gaussian closely approximates the operator that optimizes the product of *signal-to-noise ratio* and localization

J. Canny, [***A Computational Approach To Edge Detection***](#), IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

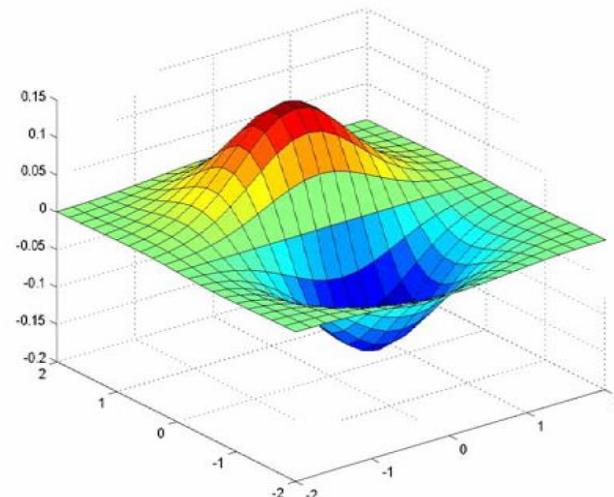


* [1 -1] =

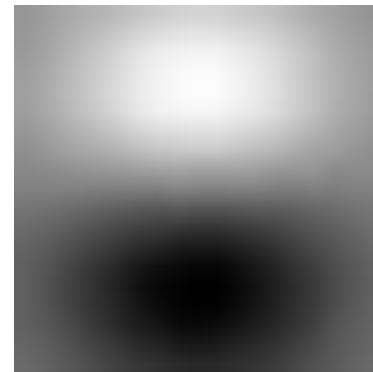
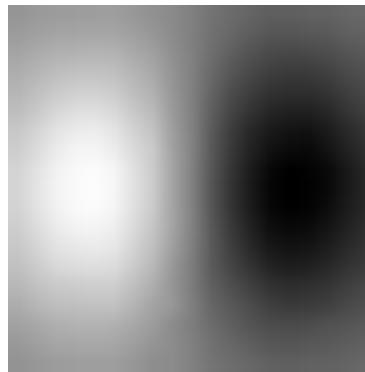




x-direction



y-direction



Example



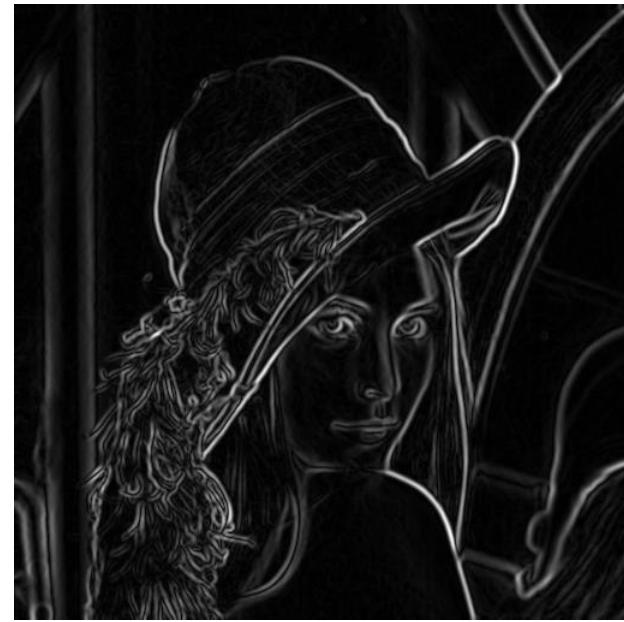
original image (Lena)



X-Derivative of Gaussian



Y-Derivative of Gaussian

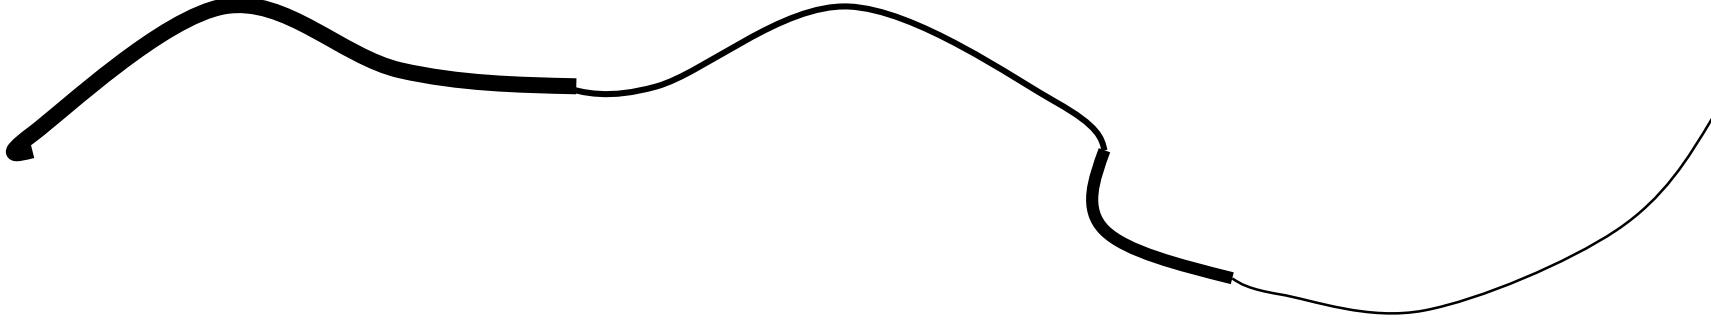


Gradient Magnitude

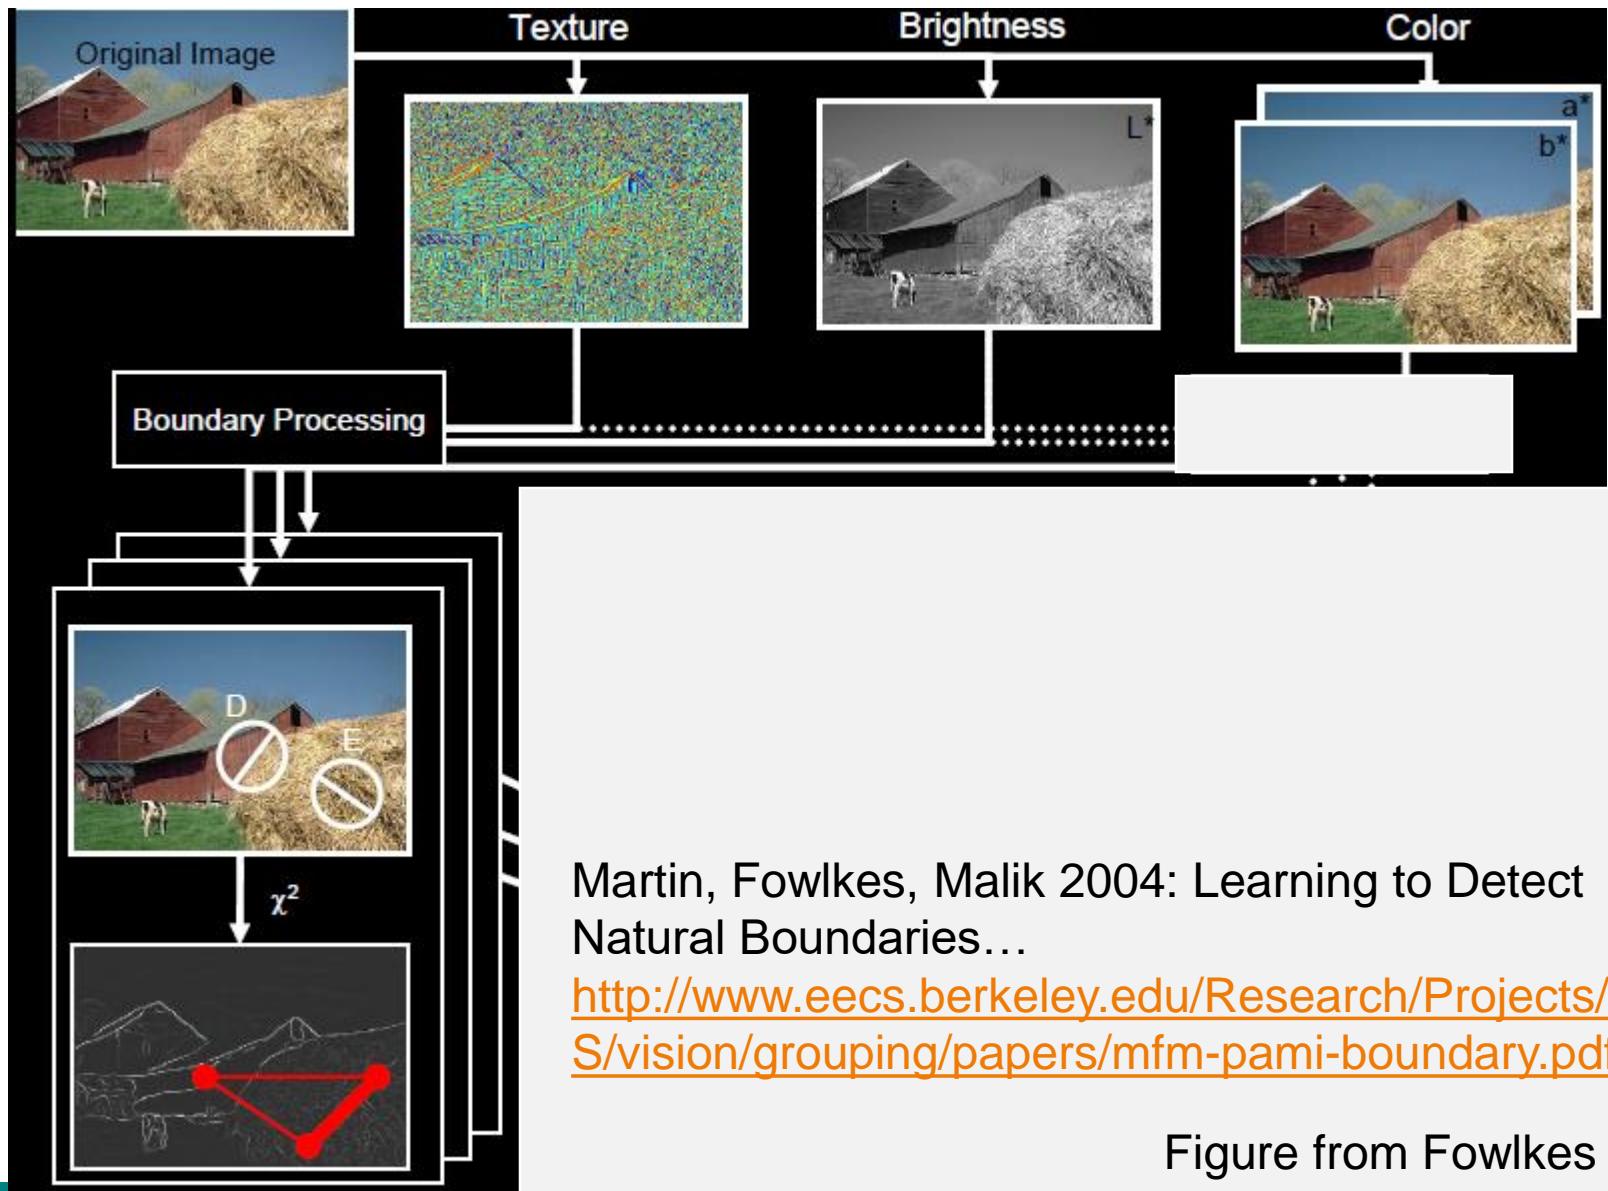
- Threshold at low/high levels to get weak/strong edge pixels
- Do connected components, starting from strong edge pixels



- Check that maximum value of gradient value is sufficiently large
 - drop-outs? use **hysteresis**
 - use a high threshold to start edge curves and a low threshold to continue them.







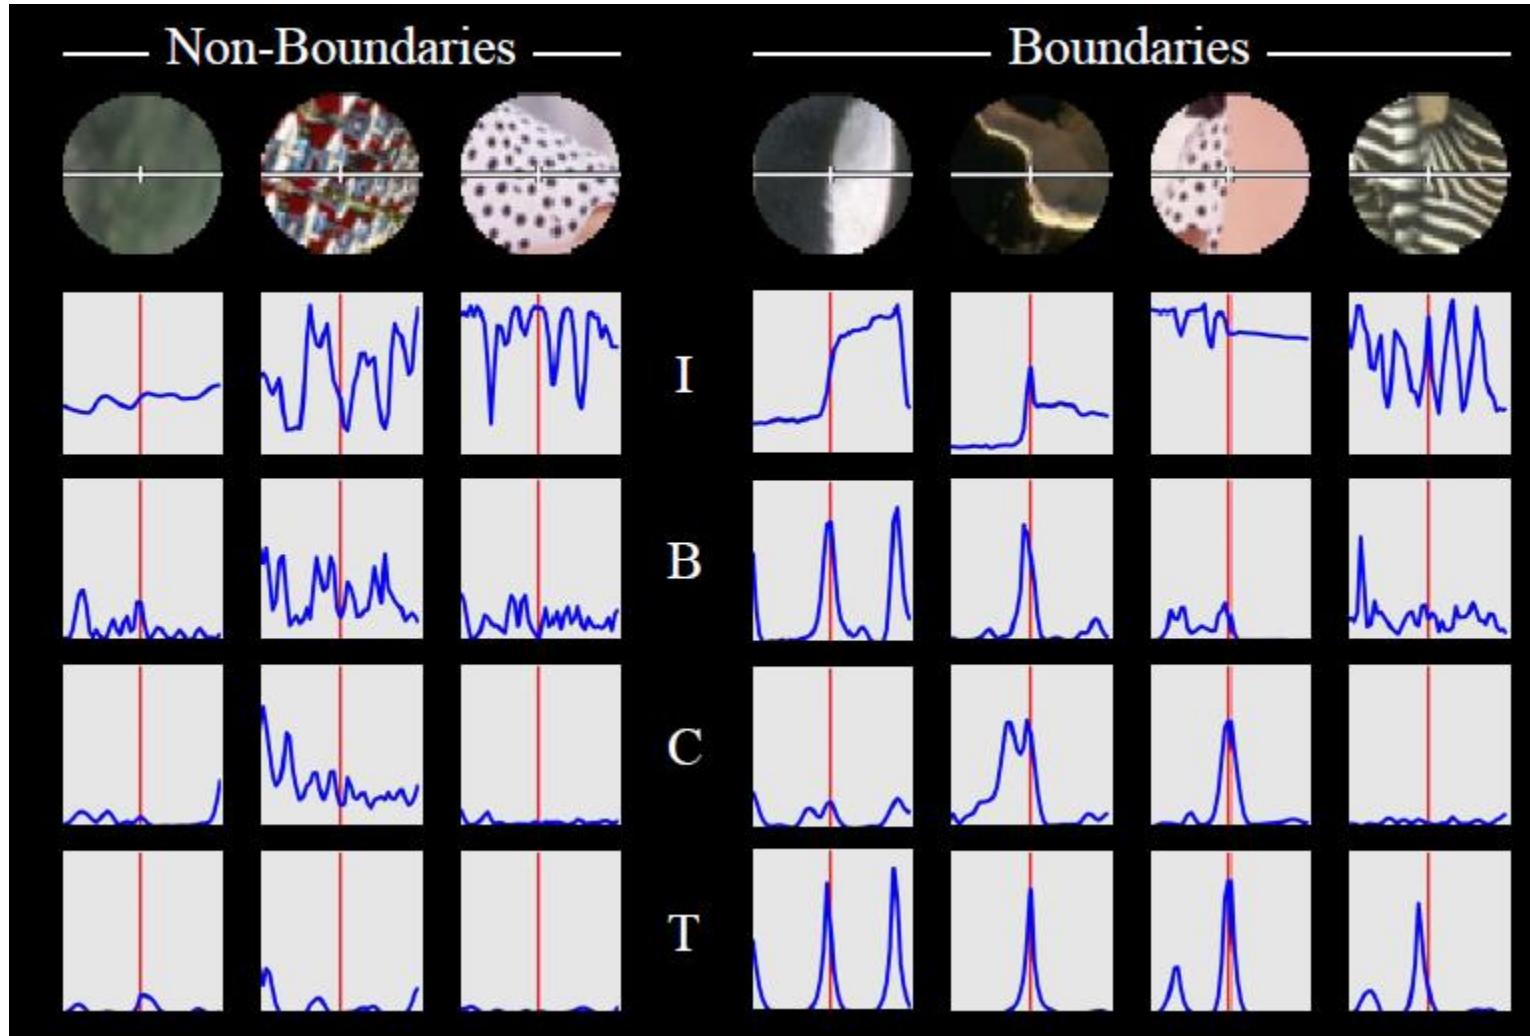


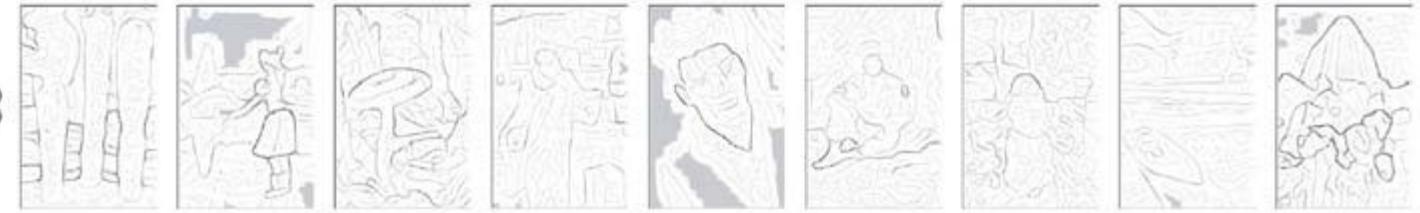
Figure from Fowlkes



Brightness



Color



Texture



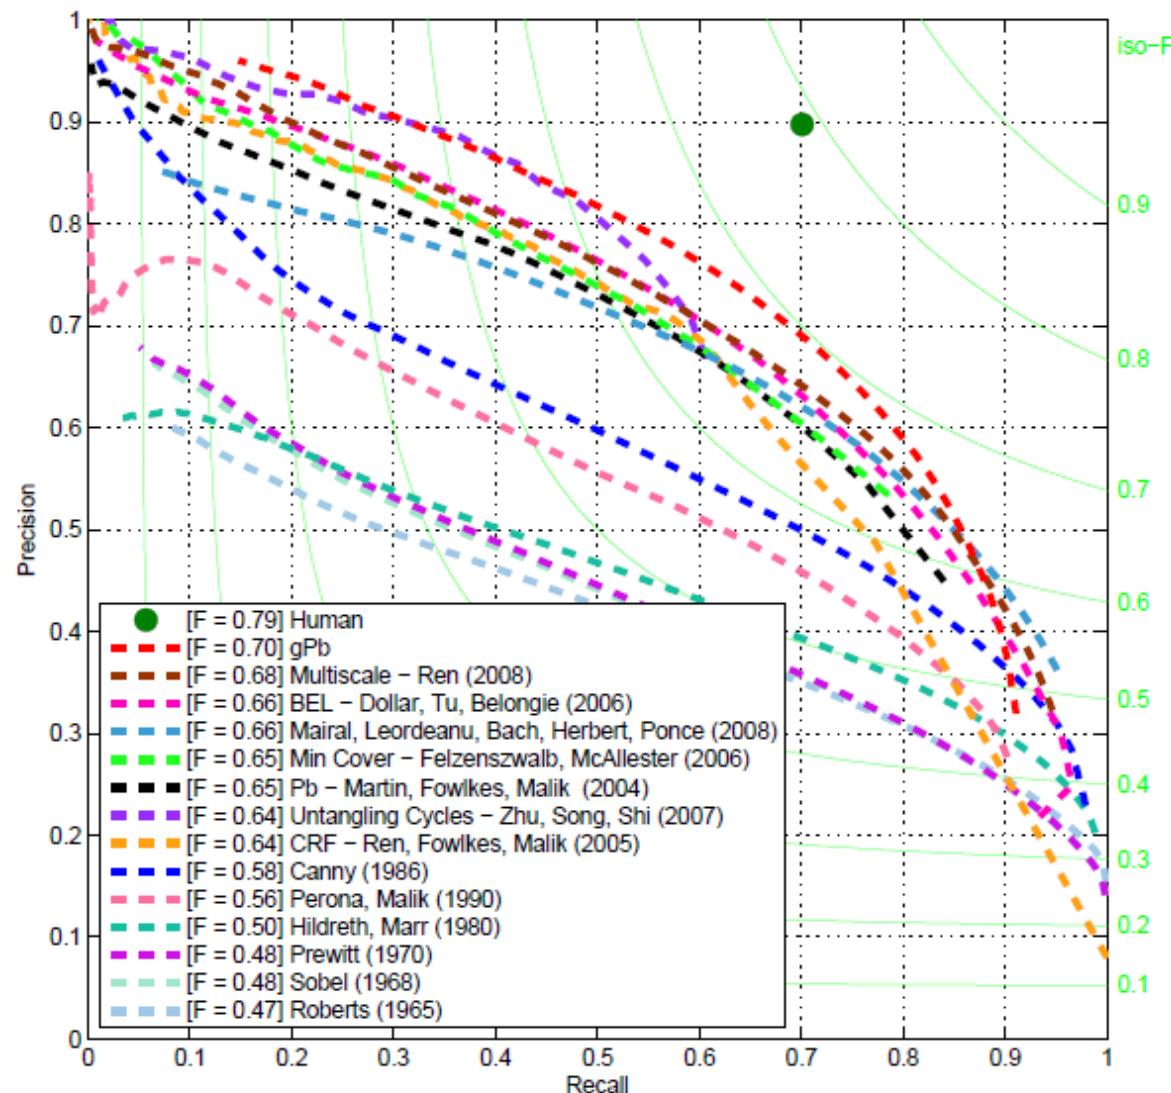
Combined



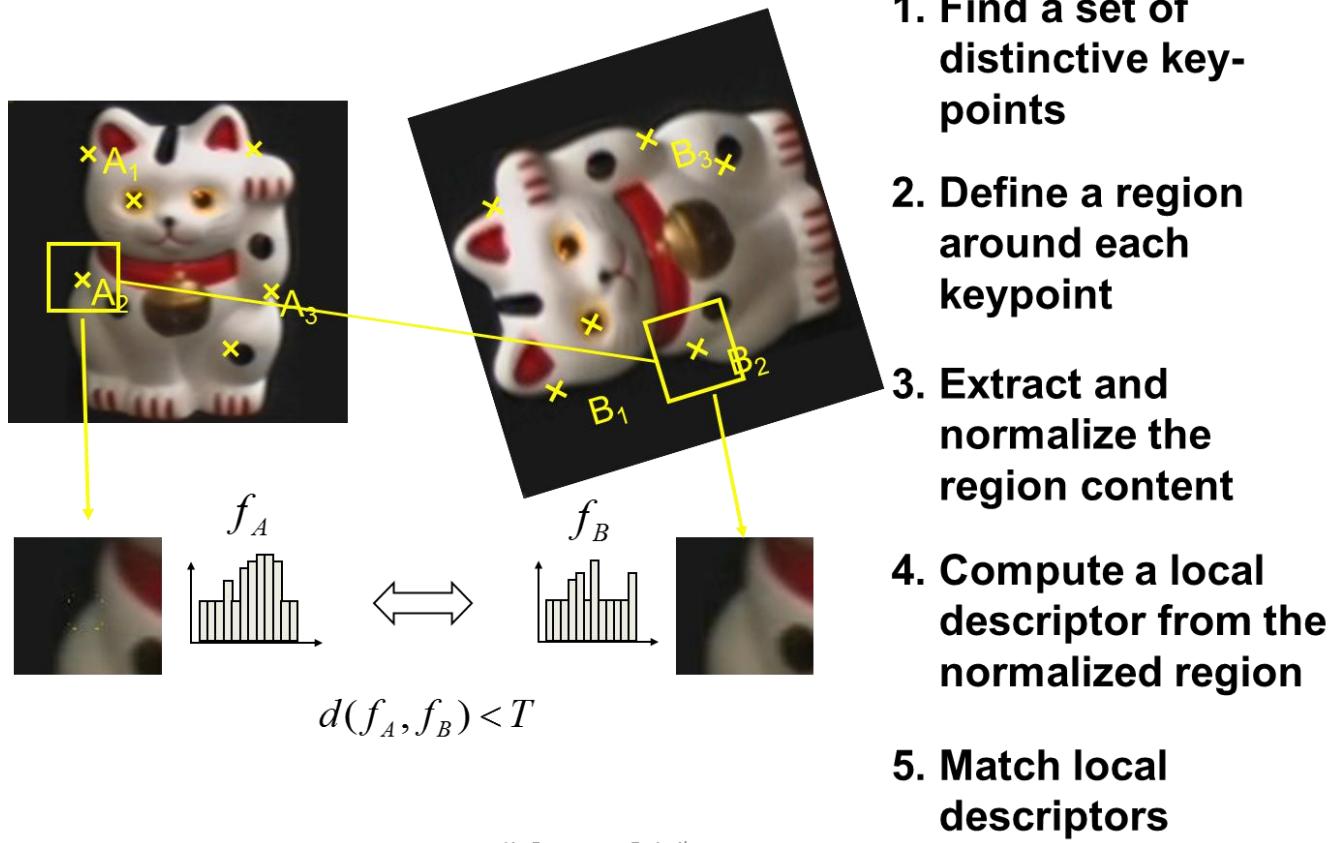
Human



45 years of boundary detection

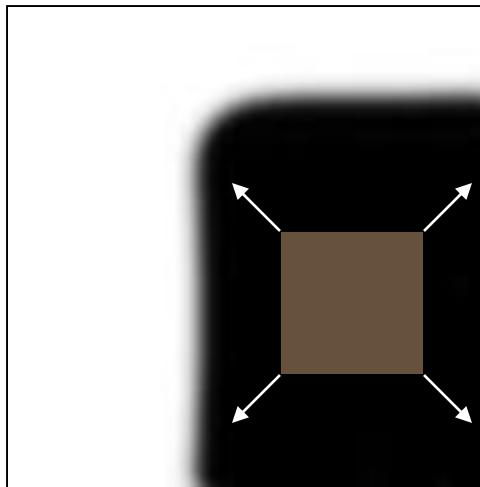


- For matching, interesting points are most concerned.
- Most of interesting points are conner:

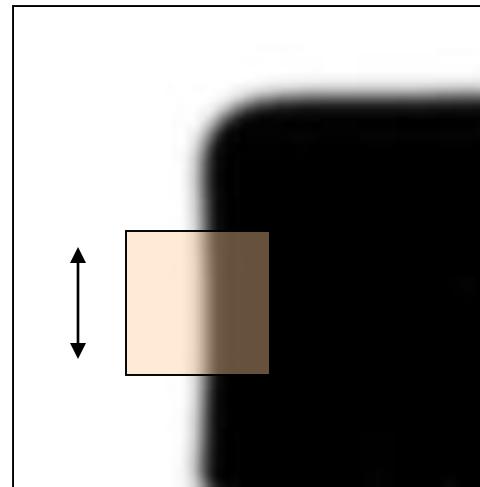


K. Grauman, B. Leibe

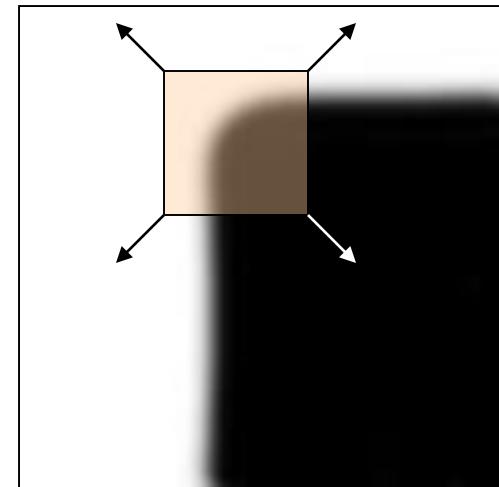
- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give *a large change* in intensity



“flat” region:
no change in
all directions



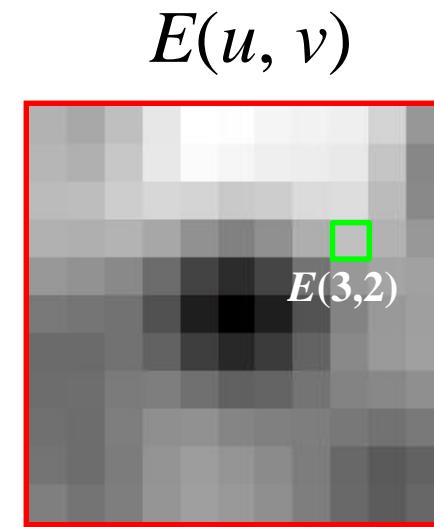
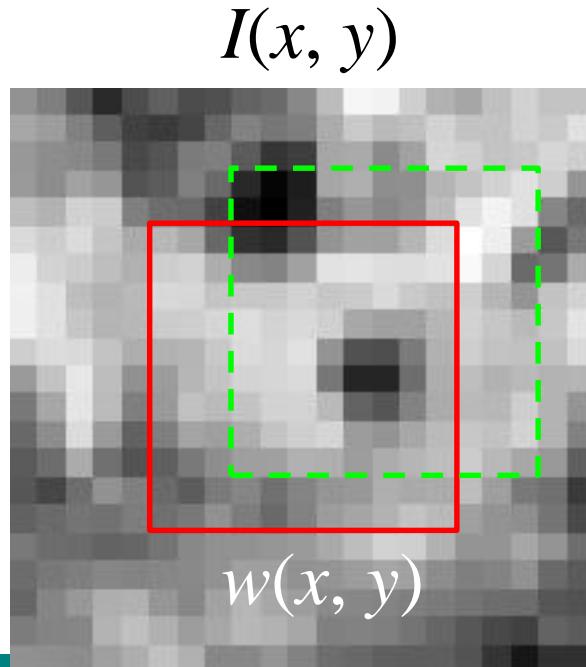
“edge”:
no change
along the edge
direction



“corner”:
significant
change in all
directions

Change in appearance of window $w(x,y)$
for the shift $[u,v]$:

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$



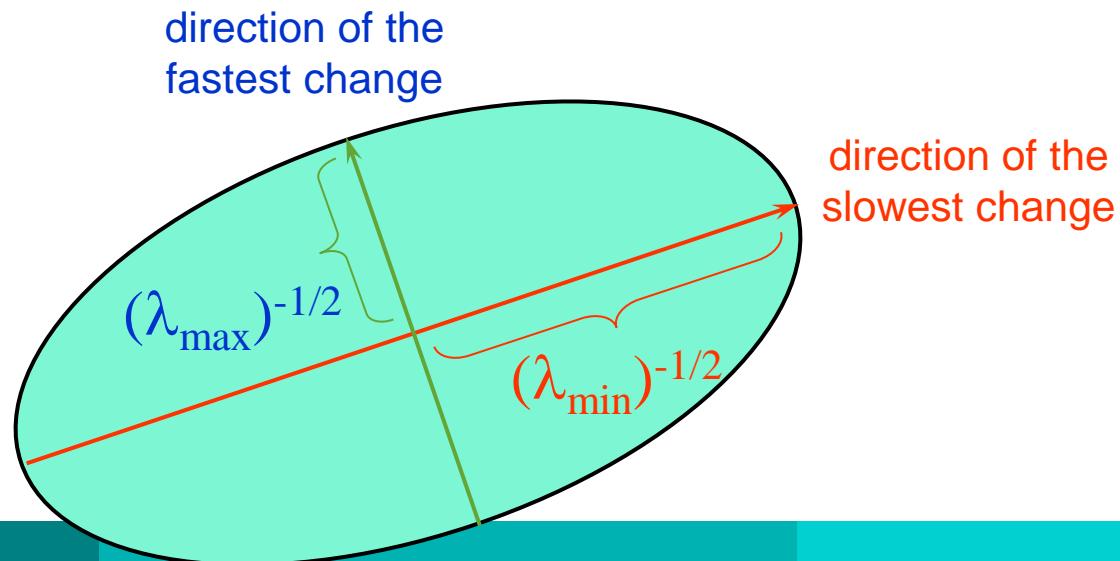
Consider a horizontal “slice” of $E(u, v)$: $[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$

This is the equation of an ellipse.

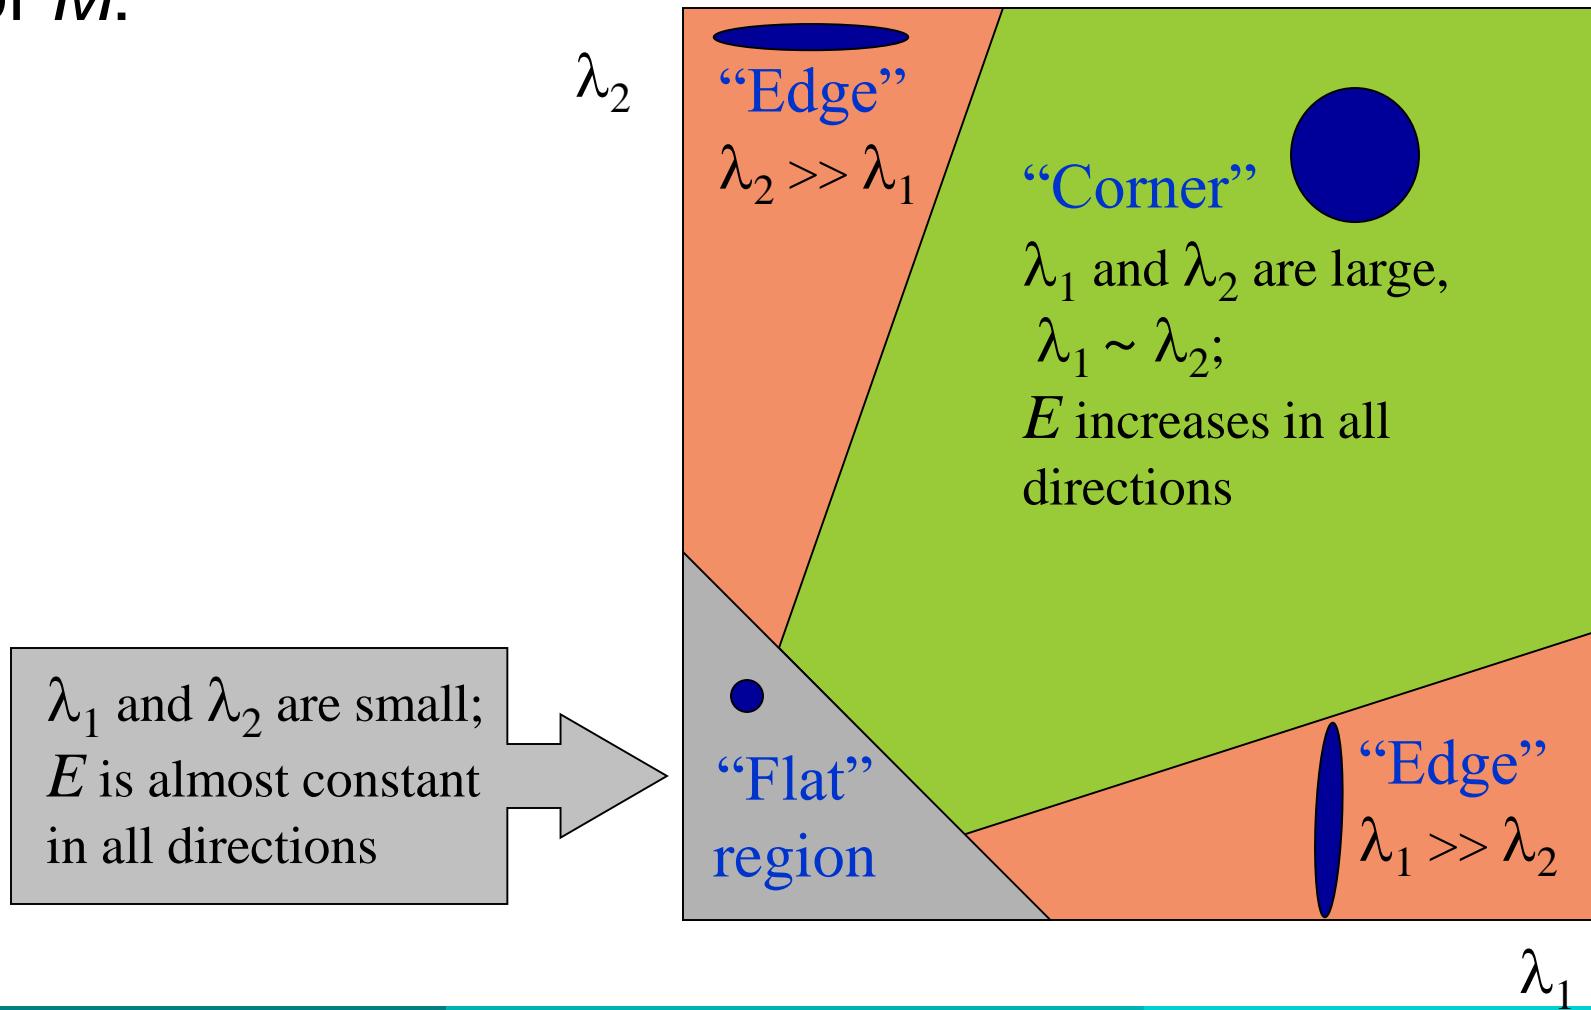
Diagonalization of M :

$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$

The axis lengths of the ellipse are determined by the eigenvalues and the orientation is determined by R

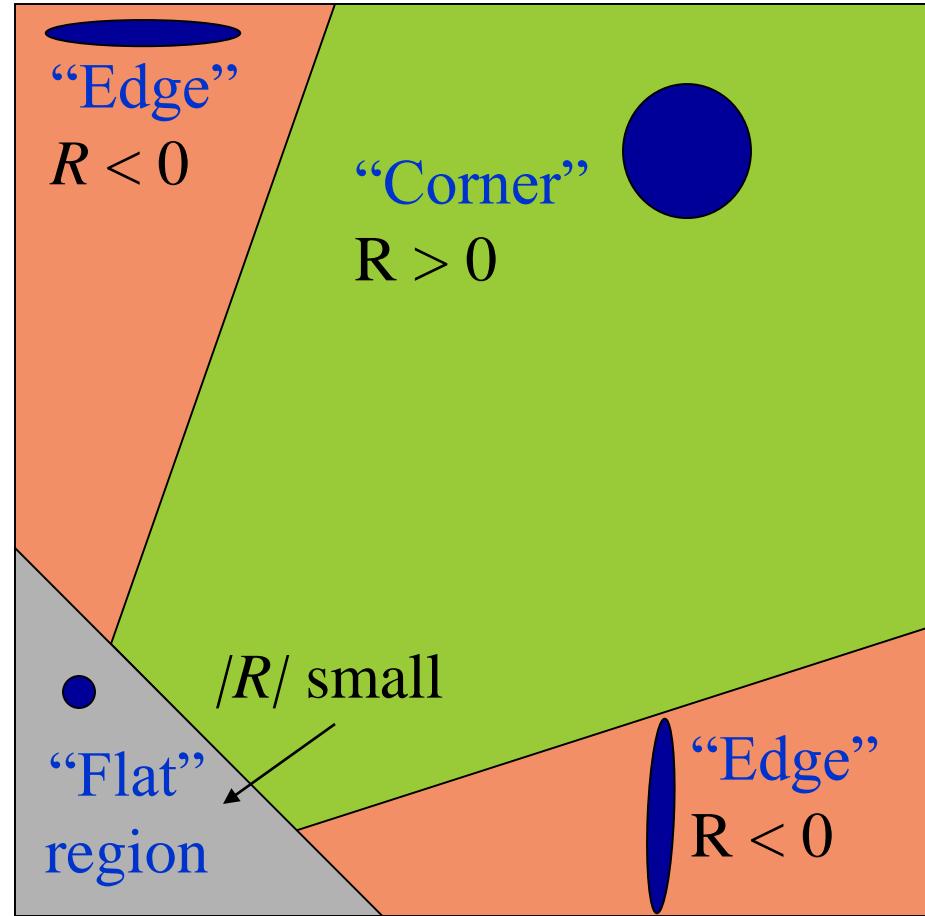


Classification of image points using eigenvalues of M :



$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

α : constant (0.04 to 0.06)





Harris corner detector

- 1) Compute M matrix for each image window to get their *cornerness* scores.
- 2) Find points whose surrounding window gave large corner response ($f >$ threshold)
- 3) Take the points of local maxima, i.e., perform non-maximum suppression

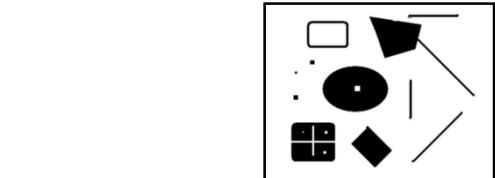
C.Harris and M.Stephens. "[A Combined Corner and Edge Detector.](#)"
Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

Harris Detector [Harris88]

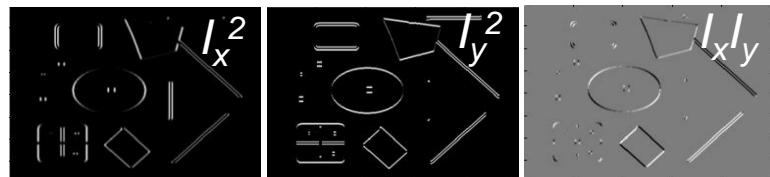
- Second moment matrix

$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

1. Image derivatives
(optionally, blur first)



2. Square of derivatives



3. Gaussian filter $g(\sigma_\nu)$



$$\det M = \lambda_1 \lambda_2$$

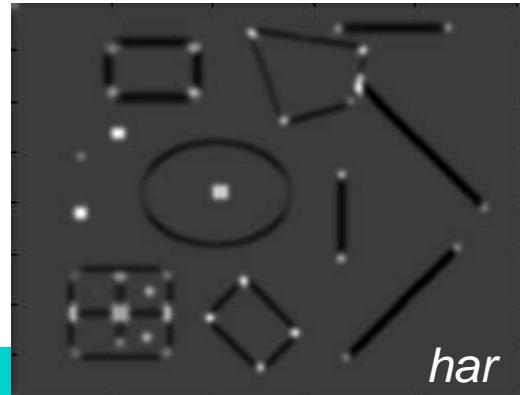
$$\text{trace } M = \lambda_1 + \lambda_2$$

4. Cornerness function – both eigenvalues are strong

$$har = \det[\mu(\sigma_I, \sigma_D)] - \alpha [\text{trace}(\mu(\sigma_I, \sigma_D))^2] =$$

$$g(I_x^2)g(I_y^2) - [g(I_x I_y)]^2 - \alpha[g(I_x^2) + g(I_y^2)]^2$$

5. Non-maxima suppression

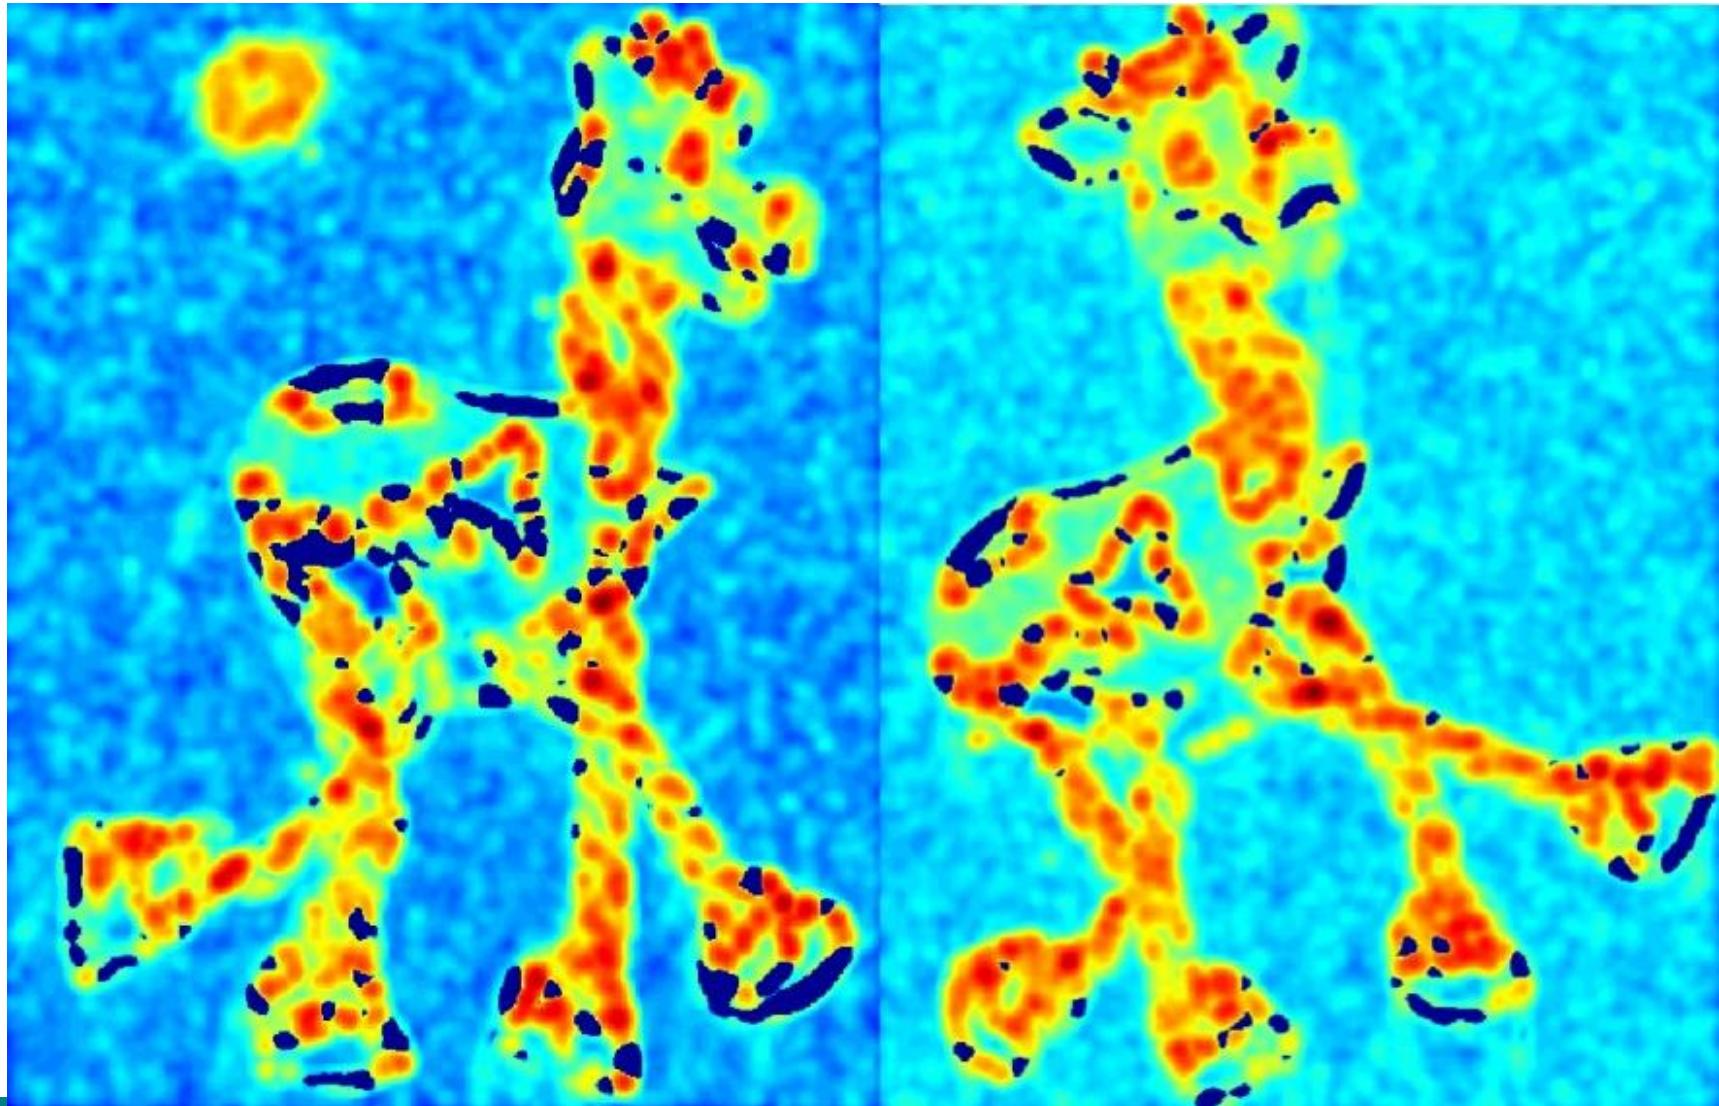




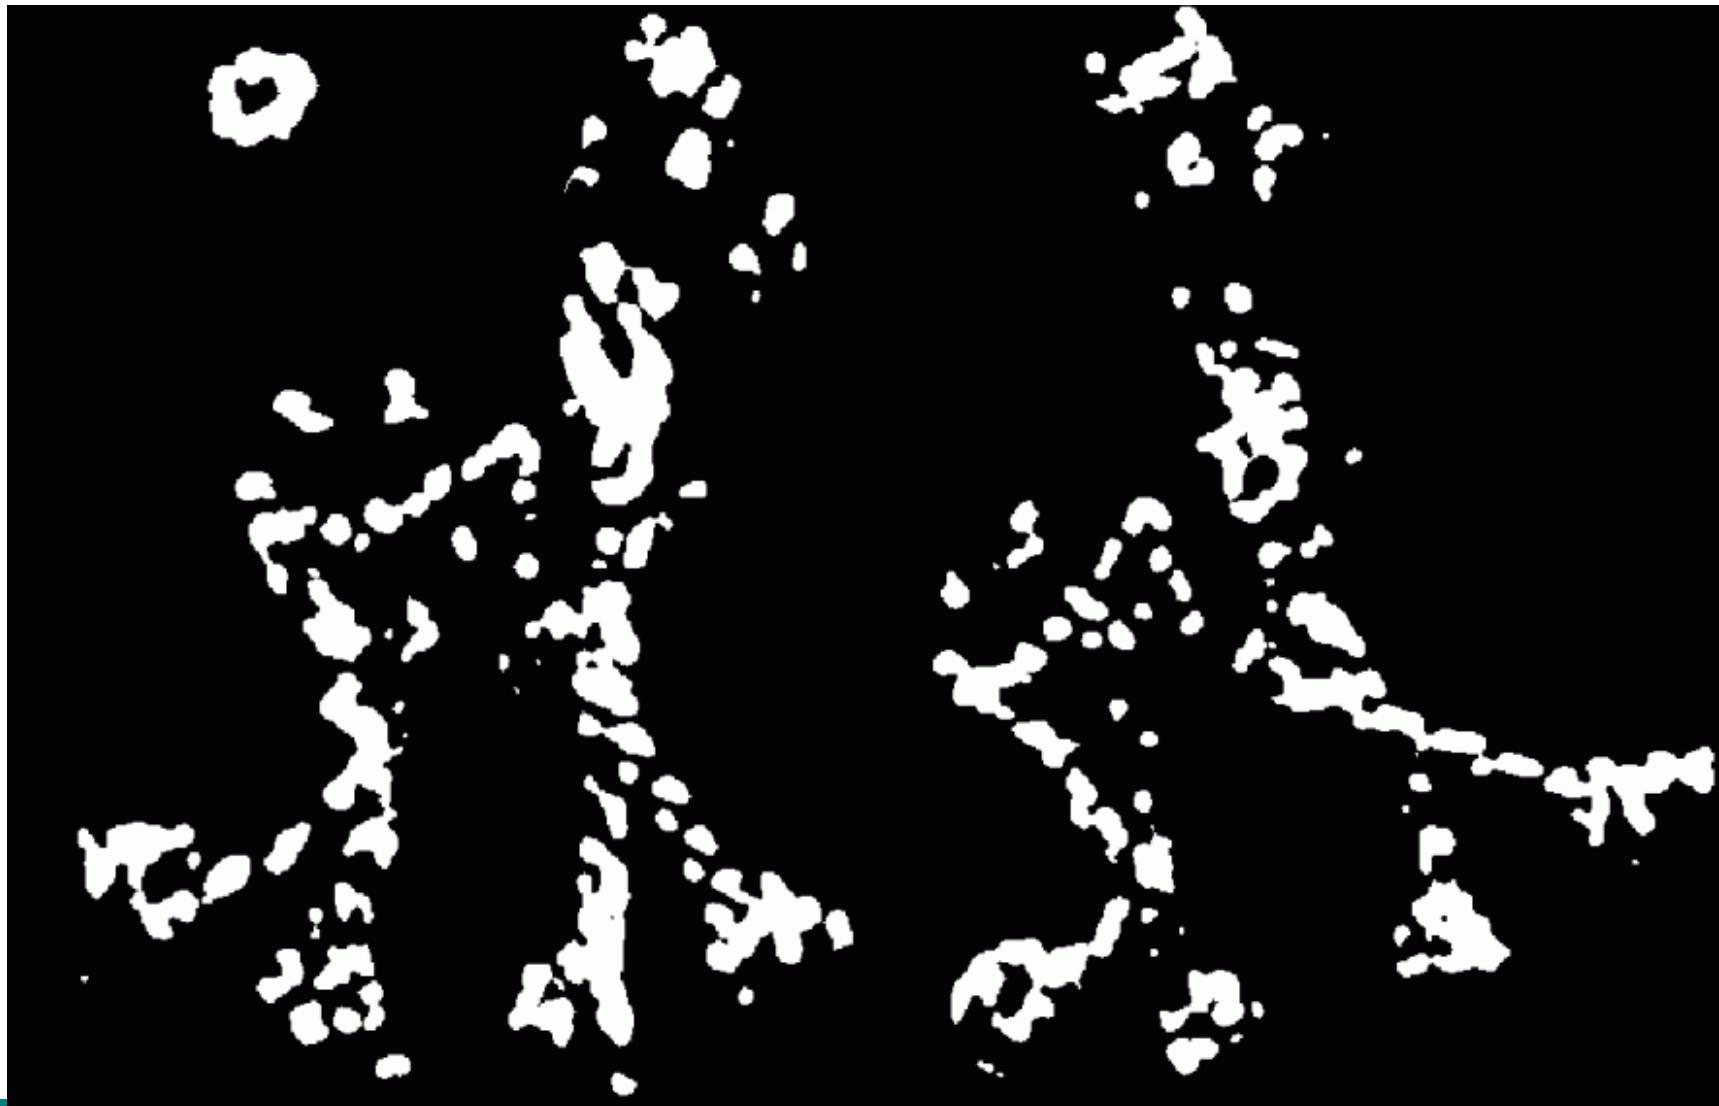
Harris Detector: Steps



Compute corner response R



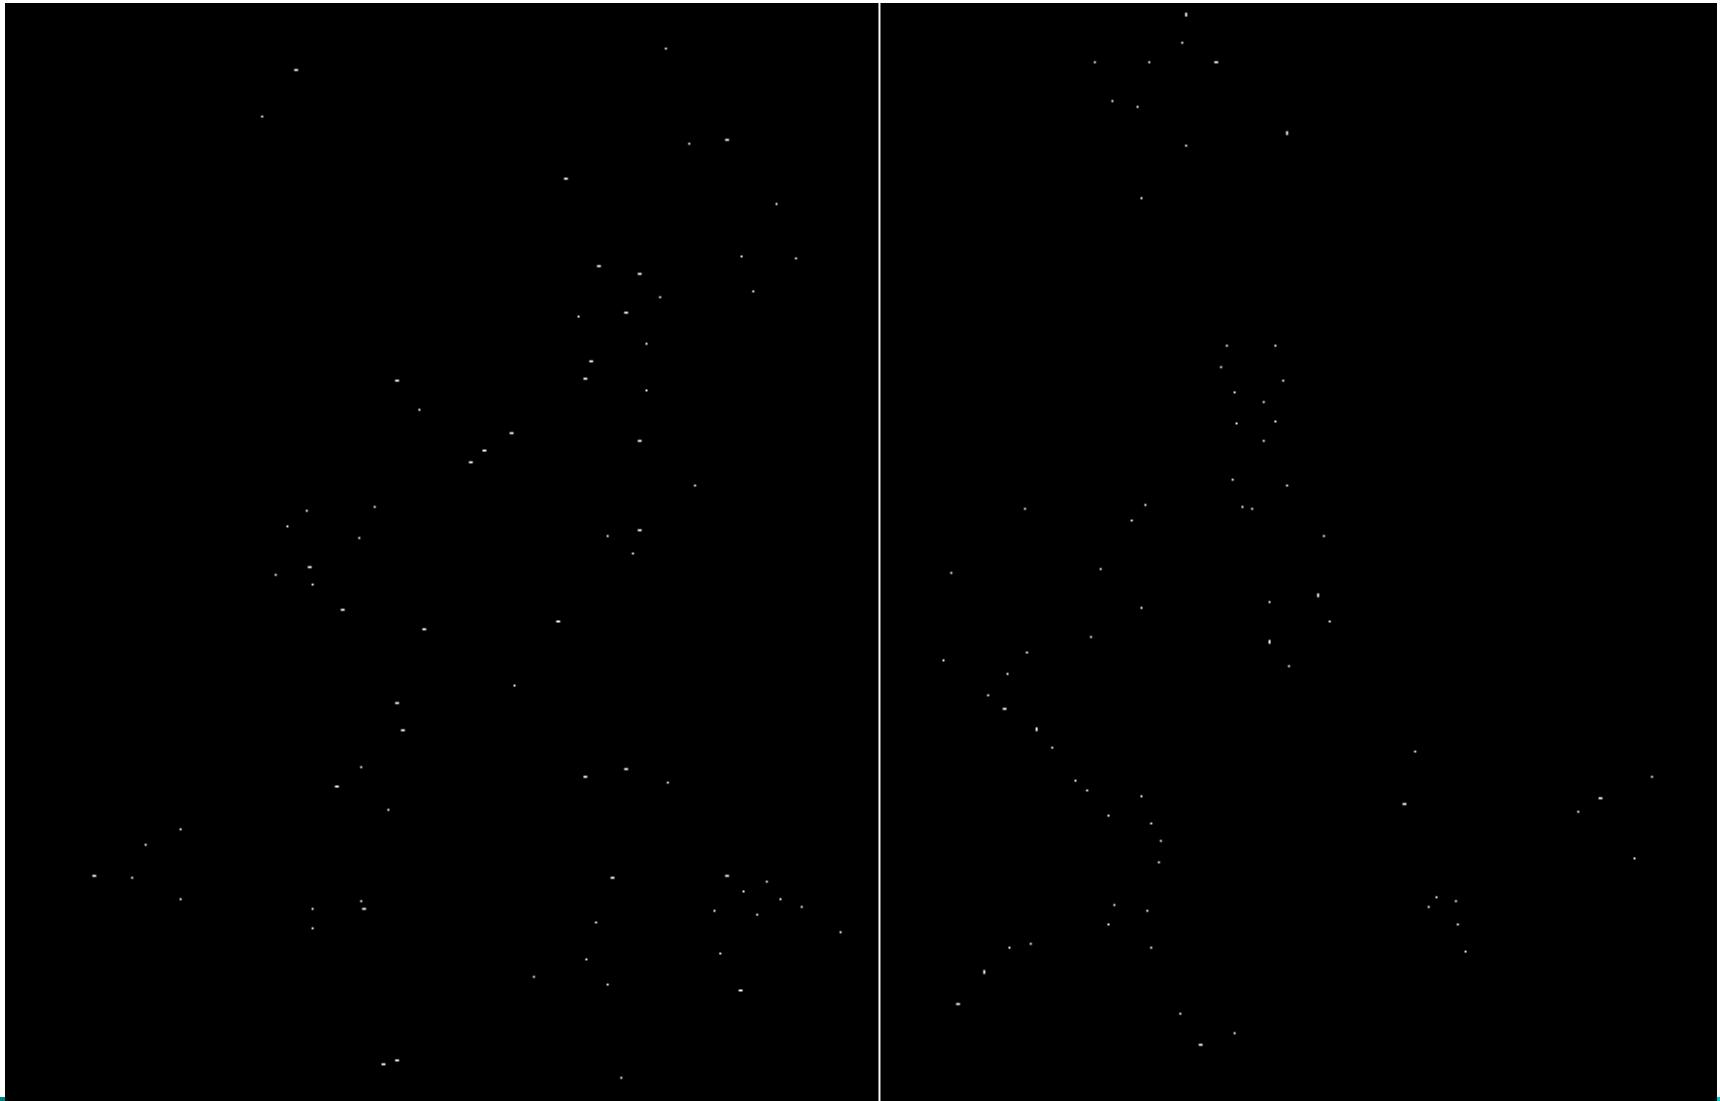
Find points with large corner response: $R > \text{threshold}$





Harris Detector: Steps

Take only the points of local maxima of R





- “Vision is the act of knowing what is where by looking.” –Aristotle
- The motivation of vision processing is let computer understand the image/video.
- To do it, the image is firstly analyzed to distract the feature points. In this presentation:
 - Basic knowledge for vision processing.
 - Filtering/Matching.
 - Border/Edge detection.
 - Conner detection.
- For more advanced processing, we need further algorithm and analysis.

- CS 143 Introduction to Computer Vision - James Hays ; <http://cs.brown.edu/courses/cs143/>
- <http://www.cse.psu.edu/~rtc12/CSE486/>
- Computer Vision: Algorithms and Applications - Richard Szeliski <http://szeliski.org/Book>

Thank you for your attention !

Let's do Question and Answer !!!